

Pr. Olivier Gruber

Full-time Professor

Université Joseph Fourier


Laboratoire d'Informatique de Grenoble

Olivier.Grubert@imag.fr

Virtual Machine Basics

2



- **Virtual Machines** versus **Real Machines**
 - *A virtual machine defines a machine (interface)*
 - *A virtual machine is a machine (implementation)*



Dec Alpha Example

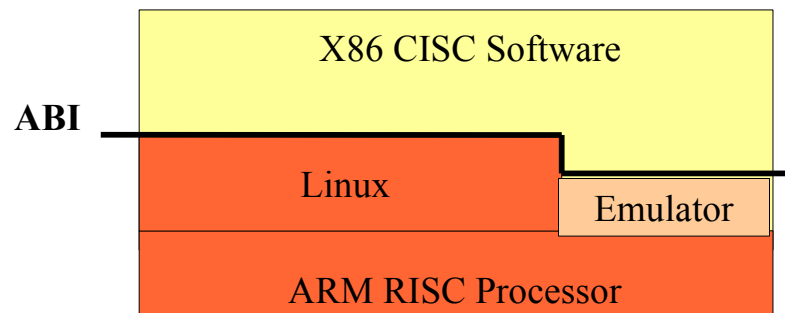
3

- Digital Alpha machine
 - Early provider of a 64bit RISC processor
 - Challenge: no existing software...
 - Support program binaries compiled to a different ISA / same ABI
 - Same ABI: ported the operating system
 - Different ISA: emulate one instruction set on a different instruction set



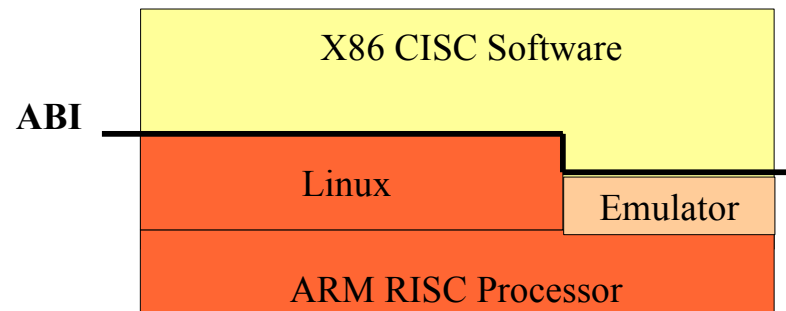
Hookway and Herdeg. *Digital FX!32: Combining Emulation and Binary Translation*.
Digital Technical Journal, January 1997, pp 3-17
Zheng and Thompson. *PA-RISC to IA-64: Transparent Execution, No Recompilation*.
IEEE Computer, March 2000, pp. 47-53

- Step One
 - *Using interpretation* (fetch, decode and emulate)
 - Design the emulator for x86 CISC to ARM RISC
 - Running inside a linux process
 - Emulating x86 code, compiled against a compatible linux ABI



- Step One Answers

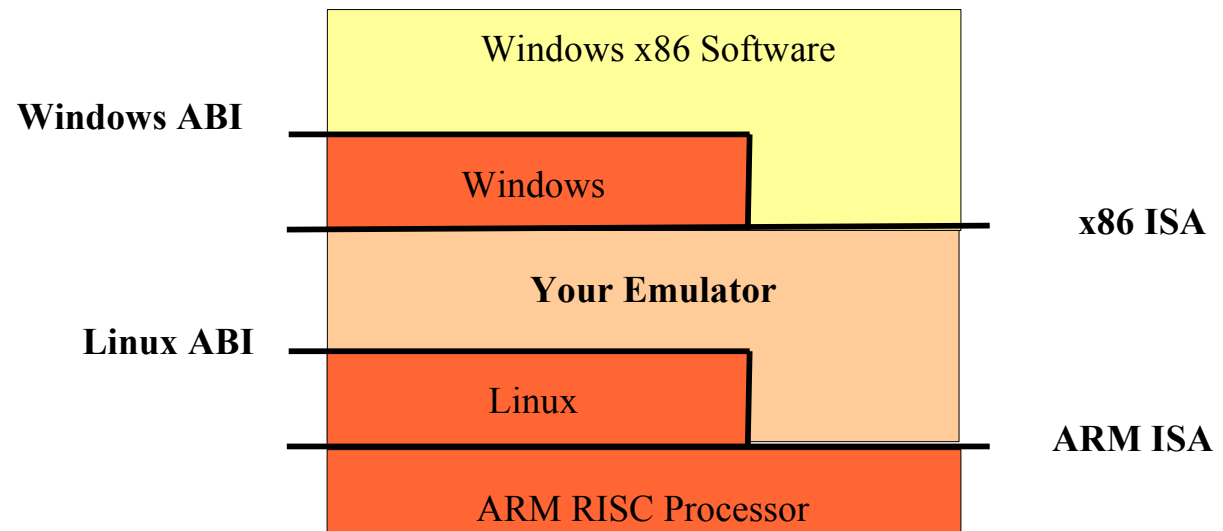
- Your emulator runs inside a process
 - Essentially about fetch, decode and emulate
- Fetching
 - Manage a program counter, being aware of branch operations
- Decode - Emulate
 - A switch like decoding leading to executing emulation snippets
 - Manage x86 registers
 - Emulate memory over a large malloc
 - Bridge the ABI calls



Design Study

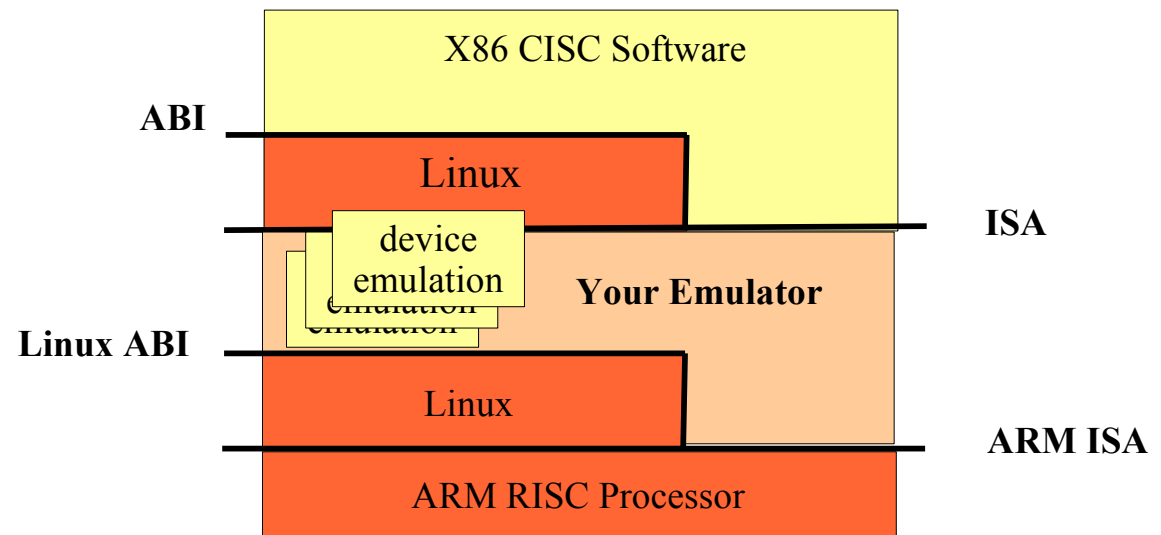
6

- Step Two
 - *Still using interpretation* (fetch, decode and emulate)
 - Design the emulator for **the complete IA-32 ISA**
 - What is different now?



- Step Two Answers

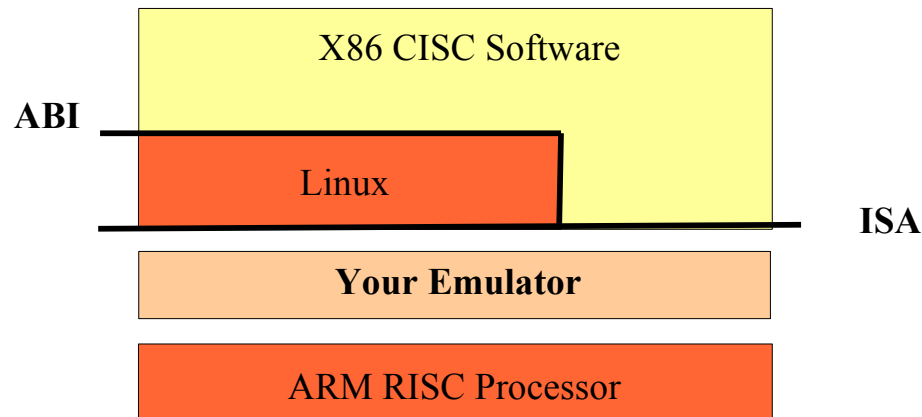
- Just a more complex architected state
- Registers and memory
- Complete interrupt/trap emulation
- Complete device emulation, leverage Linux device drivers



- **Serial lines**
 - Maybe over sockets or pipes
 - Easy byte in/out
- **Screen and Mouse**
 - Maybe over SDL
 - A window as the screen → emulate a frame buffer
 - Mouse events → encoding x,y position and buttons
- **Keyboard**
 - Scan codes / character translation
 - Assumes a keyboard layout


- Emulate a hard disk or a network card
 - Maybe over a file, or a real partition
 - Decode I/O requests
 - Automata, input from emulated hardware registers
 - Often organized as a ring buffers (in/out)
 - Issue real I/O on the file
 - Blocking I/O would be a poor emulation
 - Requires multi-threading in your emulation
 - Raise interrupts
 - Race condition with the execution
 - Interrupts only happen in between instructions

- Step Three
 - *Still using interpretation* (fetch, decode and emulate)
 - Still emulating a complete x86 illusion
 - **But a bare metal emulation** (no operating system)
 - What is different now?



- Step Three Answers


- You can no longer rely on an operating system
- You still need an IA-32 emulator
- You also needs drivers for the real devices
- And you also need an emulation for virtual devices



CISC Outside – RISC Inside

12

- Did you know?
 - Intel Processors are emulators
 - CISC code is translated, optimized, and parallelized on the fly
 - One of the key performance aspects of Intel processors
 - One of the reasons for its high energy consumption
 - Other features are possible
 - Speculative execution for example
 - ***So real machines are not so real after all...***



Virtual Machine – So what are they?

13

- Operating Systems


- Unix – Linux – Windows, almost a 40 year-old design
- Different approaches exist(ed)
 - Palm-OS – a refreshing approach
 - AS-400 – a different world
 - Mac-OS with micro-kernel technology

- Hypervisors

- New or old technology?
- Both of course...

- High-level languages – a different future?

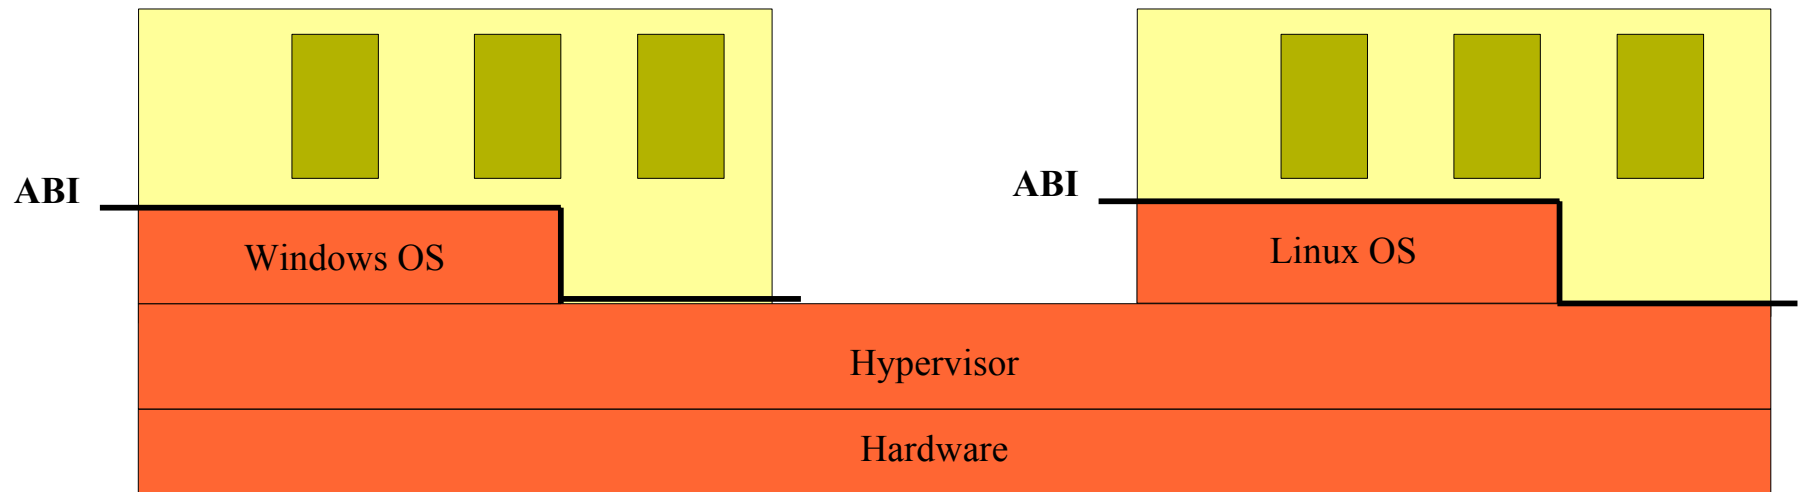
- Microsoft DotNet – quite a fundamental change
- Java bare metal or Arduino
- Google Android



Hypervisor – Basics

14



- Bare-metal virtualization (Type-I or System VMs)
 - Such as Vmware ESX or Xen or Nova
 - Multiplex operating systems, often para-virtualized
 - Often virtualize a similar hardware (but not always)
 - Often relies on hardware-assisted virtualization



Hypervisor – Basics


15

- In-Process Hypervisor (Type II or Process VMs)
 - Such as VirtualBox or QEMU
 - Virtualizing a real machine within a process
 - Para-virtualized or out-of-the-box operating systems
 - Same hardware (QEMU/VirtualBox) or not (QEMU)



Hypervisor – Overview


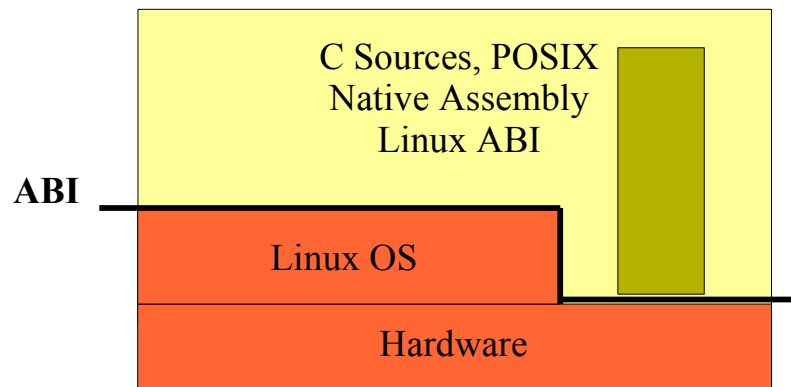
16



High-Level Virtual Machine Basics

17


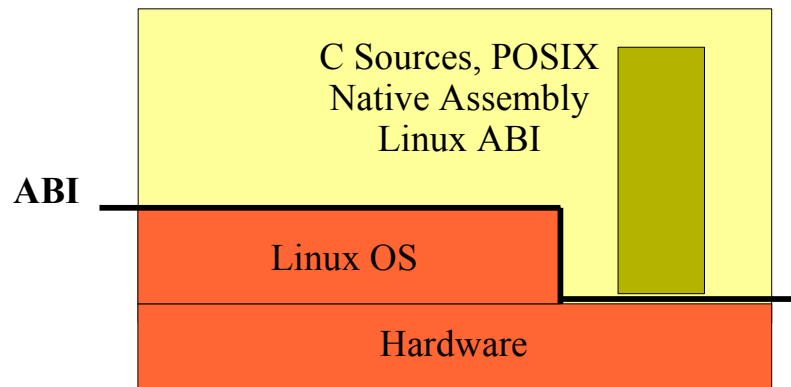
- High-level language virtual machines
 - Many are just about yet another language
 - Such as Python, Perl, JavaScript, or Java
- Virtualize
 - Provide a different instruction set
 - Sometimes provide different libraries



High-Level Virtual Machine Basics

18


- But some are full-fledge platforms
 - Microsoft C# and DotNet platform
 - Eclipse Rich Client Platform, based on OSGi
 - JONAS Web Server, based on OSGi, a servlet and bean containers



High-Level Virtual Machine Basics

19

- Towards language independence
 - Microsoft Common Language Infrastructure
 - Common Language Runtime (CLR) and Common Type System (CTS)
 - The Java Virtual Machine is going in the same direction
 - Already a target for many languages (JavaScript, Scheme, Perl, Python, etc.)
 - LLVM Project
 - Interesting path




Compilation & Linking

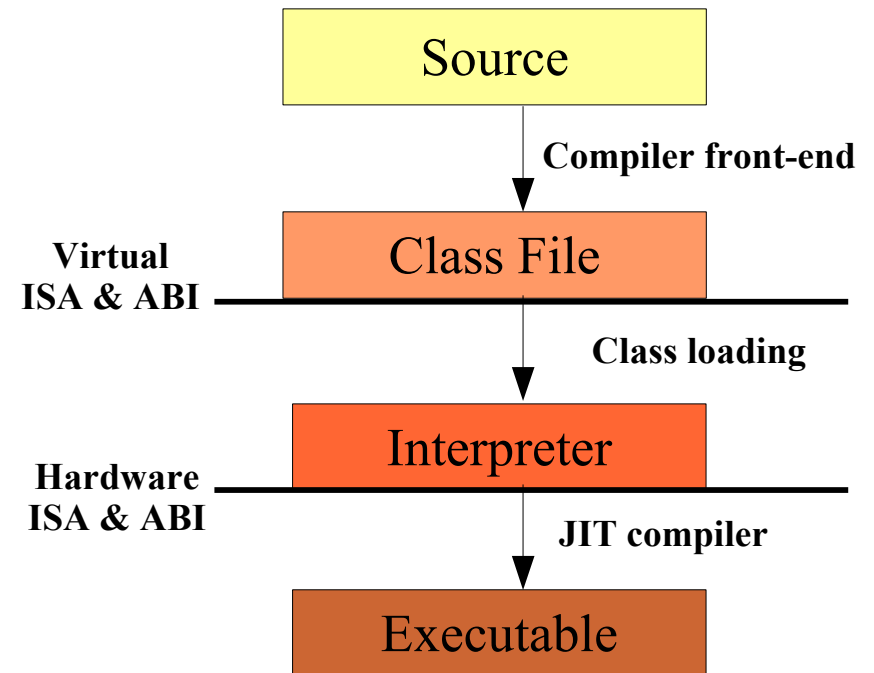
20

- A shift of responsibilities...

Traditional Language Compilation & Linking Chain



High-Level Language Compilation & Linking Chain



Virtual Machine – Everywhere

21

