

Micro-Kernels

Pr. Olivier Gruber

Université Joseph Fourier

Laboratoire d'Informatique de Grenoble

Olivier.Gruber@imag.fr

Acknowledgements and References

- This lecture includes material courtesy of:
 - Dr. Renaud Lachaize, UJF
 - Prof. Gernot Heiser, UNSW
- References:
 - Research publications cited in the slides
- Reference Books:
 - A. Tanenbaum, Modern Operating Systems (3rd Edition), Prentice Hall, 2007
 - A. Silberschatz et al, Operating System Concepts (9th Edition), Wiley, 2013

- Traditional kernels
 - Windows or Linux
 - Kernel is a single large binary, executing in kernel mode
 - Usually load/unload modules such as drivers or file systems
- Pros
 - Considered trusted – have been around and used for a long time
 - Performance – tight integration, kernel mode
 - Extensibility – if you can dream it, you can implement it

- Cons
 - Large code base
 - Millions of lines of code, a large part is driver code
 - Linux: 70% of the code is driver code
 - Hard to understand, maintain, tune, or shrink
 - Complex and therefore buggy (5-10 bugs per KLOC)
 - Bugs often cause a complete system failure (panic – blue screen)
 - Linux: 70% of kernel failures are due to drivers
 - Windows: 85% of failures are attributed to drivers

- Efforts to fix the driver problem
 - Many research efforts with tools to improve driver trustworthiness
 - Revisiting driver framework and security mechanisms
- Efforts to tackle complexity through modularity
 - Relatively small, self-contained components
 - Well-defined interfaces
- Does not work well with monolithic kernels
 - Because all kernel code executes in privileged mode
 - Interfaces cannot be enforced
 - Faults cannot be contained
 - And performance ultimately takes priority over modularity

- Academic study of the Linux kernel evolutions (2002)
 - Looked at size and coupling of kernel “modules”
 - Coupling: interdependency via global variables
 - Result 1: Module size grows linearly with version number
 - Result 2: Interdependency grows exponentially with version number!
- Intel study on Linux kernel releases (2009)
 - 12% performance drop over the last ten releases
 - Quoting Linus Torvalds (creator of Linux):
 - “We are getting bloated and huge. Yes, it’s a problem.”
 - “Our instruction cache footprint is scary. [...] And whenever we add a new feature, it only gets worse.”
- What about other kernels?
 - There are no reasons to believe that other kernels are different

Introducing Microkernels

7


- Core idea
 - Improve kernel maintainability and trustworthiness
 - Minimizing the Trusted Code Base (TCB) running in privileged mode
 - Seminal ideas can be traced to Brinch Hansen's Nucleus (1970)
- Main principles
 - Small (privileged) kernel providing core functionality
 - Most OS services provided by user-level servers
 - Servers communicate via messages

Operating System Architecture

8

- **Microkernels**

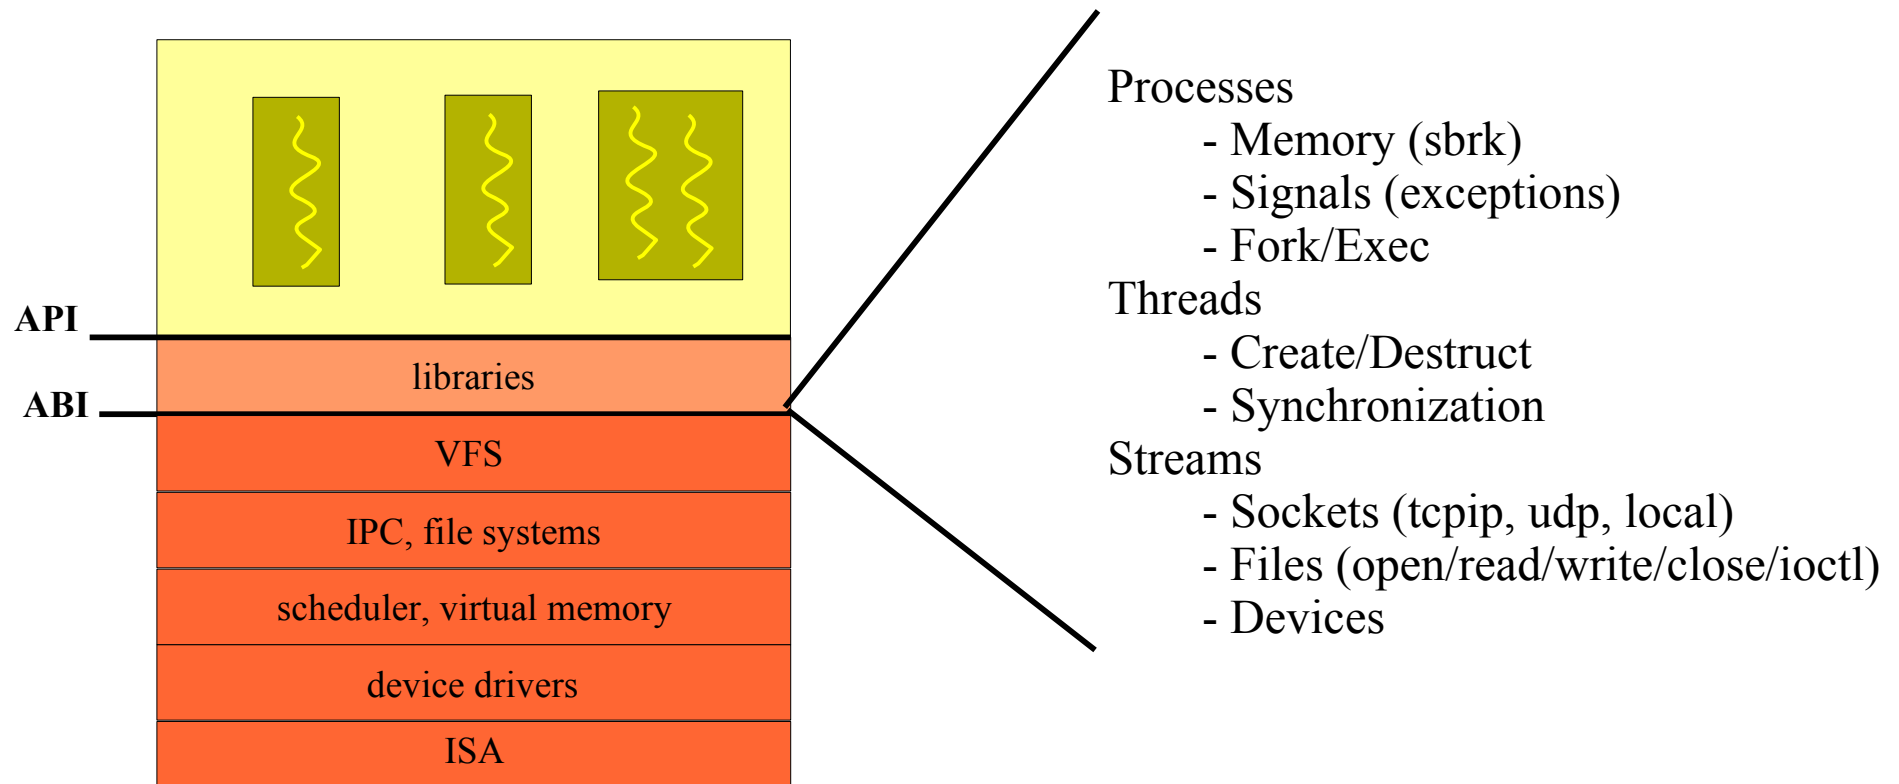
- Minimal kernel (TCB) with only few and low-level mechanisms
- Promote a modular design of cooperative user-level servers
- Cooperation relies on Inter-Process Communications (IPC)



Operating System Architecture

9


- Linux/Unix Kernel Architecture
 - The pieces and the rationales...



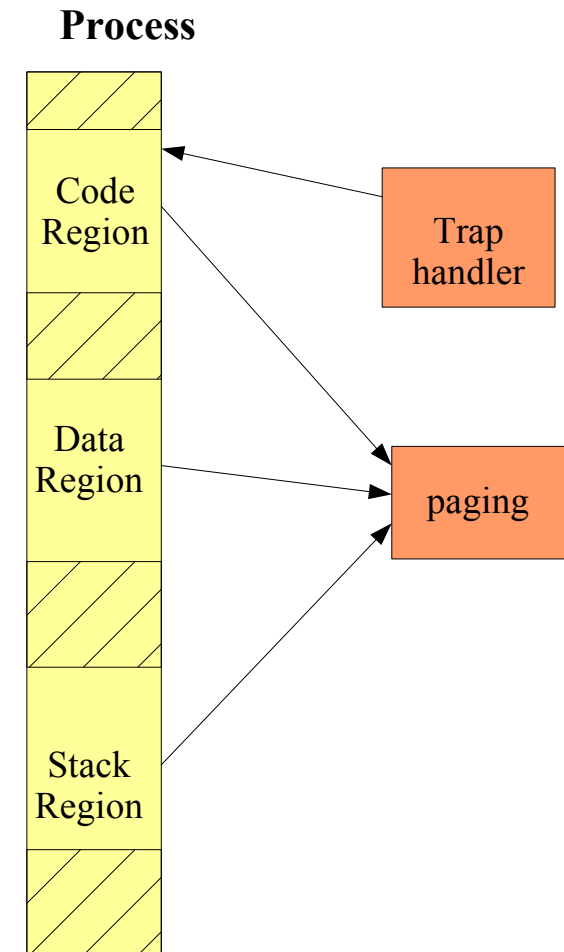
Processes

10

- Processes as address spaces
 - A virtualized memory, isolated by hardware
 - Trap handler generates signals
- Mapped memory regions
 - Paging memory regions
- Management
 - Create by forking (duplicate process)
 - Kill by pid
 - Exec (loader like Elf loader or script loader)



Paging and loaders require
file system support
in the Kernel



Boot Process

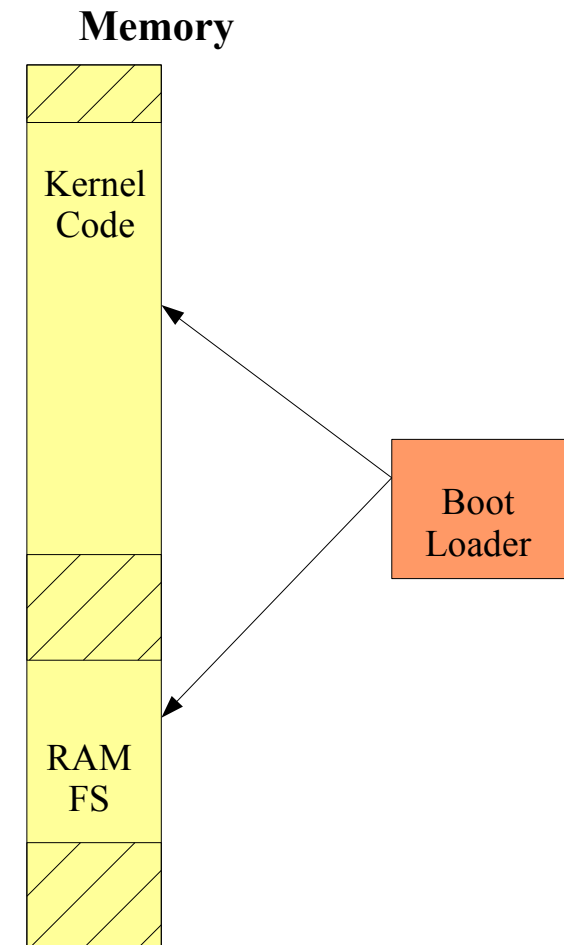
11

- First process
 - Traditionally /init
 - From the root partition /
 - Once the kernel has booted
- Boot loader
 - Loads the kernel in memory
 - From /boot

➡ Boot loader requires file system support as well

➡ BIOS loads the boot loader
So it needs file system support


➡ And therefore everyone needs device drivers



Network and Thread Support

12


- Before the early 80's
 - No network support
 - No kernel threads
- Today's kernels
 - Light-weight processes
 - Known to the scheduler
 - Generic scheduling
 - Full network support



What are the rationales
for kernel support?

- Green threads

- Can be entirely user-level library
- Fast synchronization
- Application-specific scheduling is possible
- Requires cooperative scheduling
- Blocking or non-blocking I/Os



Requires non-blocking I/Os
from the kernel...

- Kernel threads


- Light-weight processes
- Known to the scheduler
- Generic scheduling

And what about paging?

Network and Thread Support

14

- Network stack
 - Device drivers for various NICs
 - Full IP – ARP – ICMP – UDP – TCPIP support
 - Supports socket streams
- Kernel support
 - Sockest are known to the Clib API
 - Sockets are known to the ABI



Why?

Why not just a block device?


Graphic Support

15

- X11 Server
 - User-space graphic server
 - User-space device drivers
- Basic kernel support
 - Basic generic device drivers
 - Like a framebuffer originally
 - Early graphics ~ gray-scale pixels

video memory → mapped file
mode change → ioctl

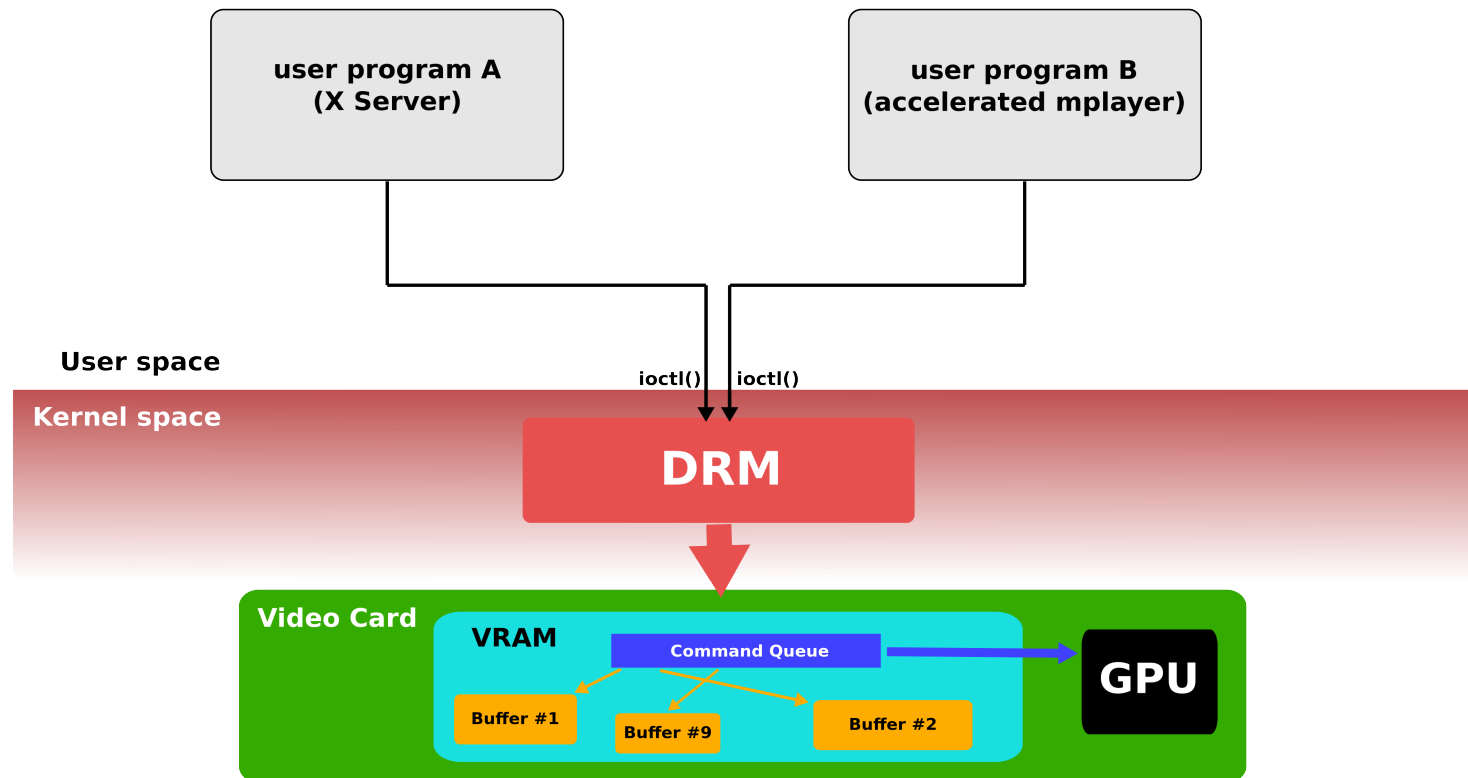
2D Acceleration?
bitblit → ioctl



Graphic Support

16

- Direct Rendering Infrastructure (DRI)
 - Direct Rendering Manager

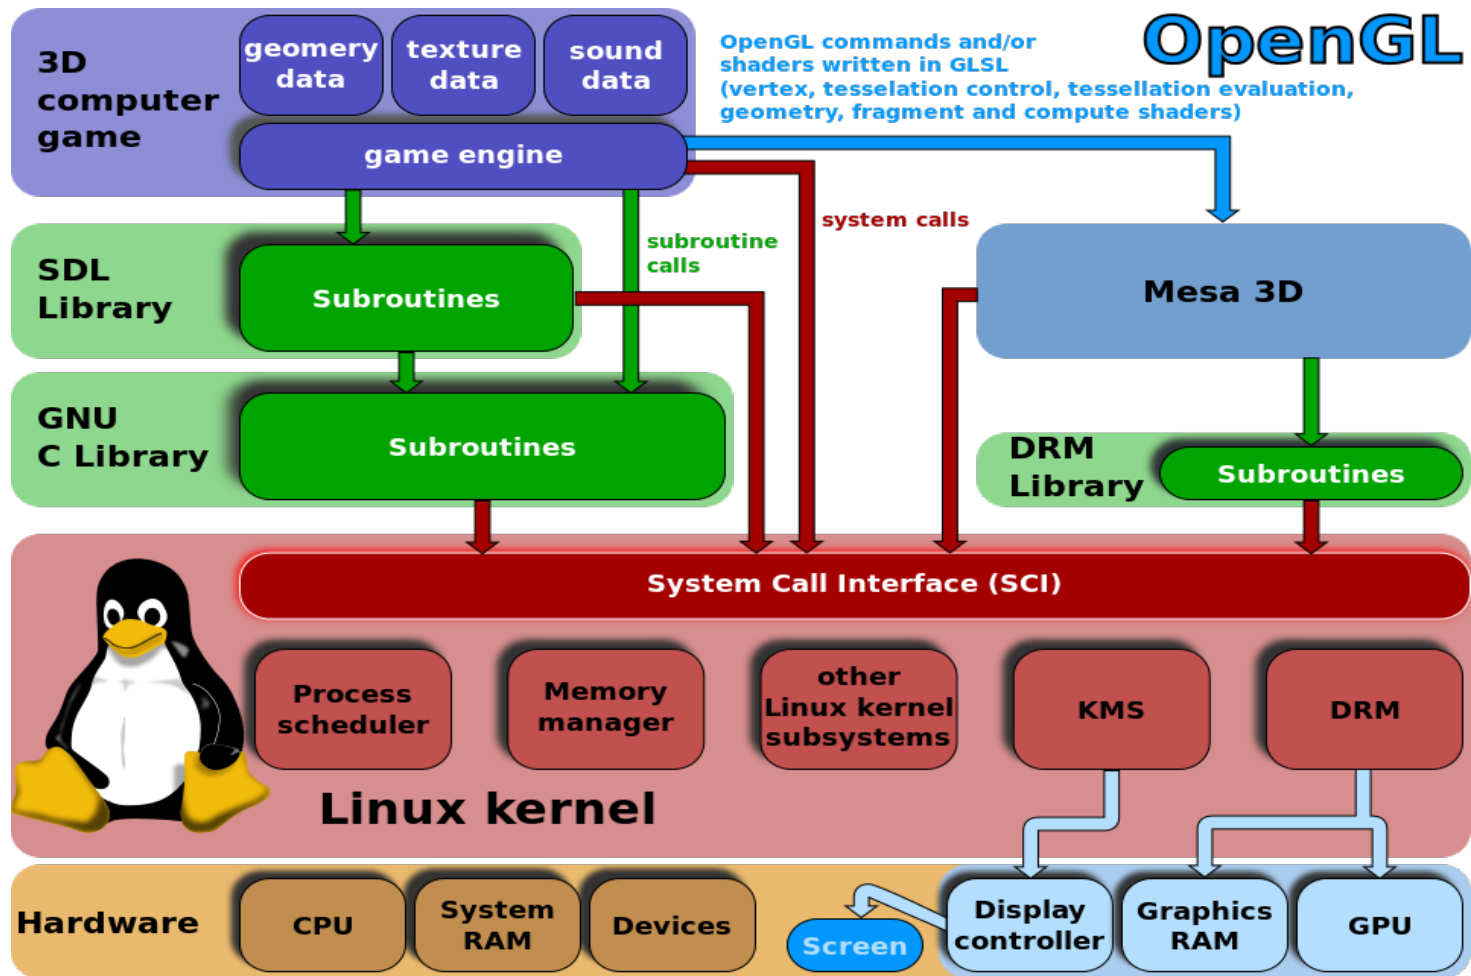


© 2014 Javier Cantero - this work is under the Creative Commons Attribution ShareAlike 4.0 license

Graphic Support

17

- Direct Rendering Infrastructure (DRI)

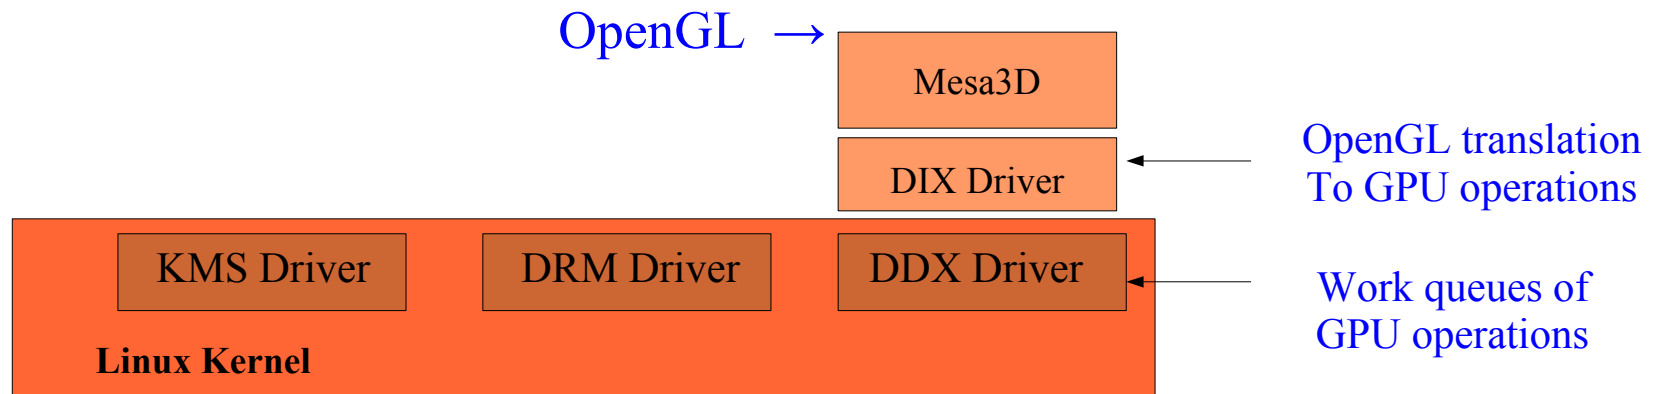


Linux kernel and OpenGL video gamesCC BY-SA 3.0
ScotXW - Own work This image includes elements
that have been taken or adapted from this: Tux-shaded.svg .

Graphic Support

18

- Mesa3D
 - Open-source OpenGL
 - Simplified overview...



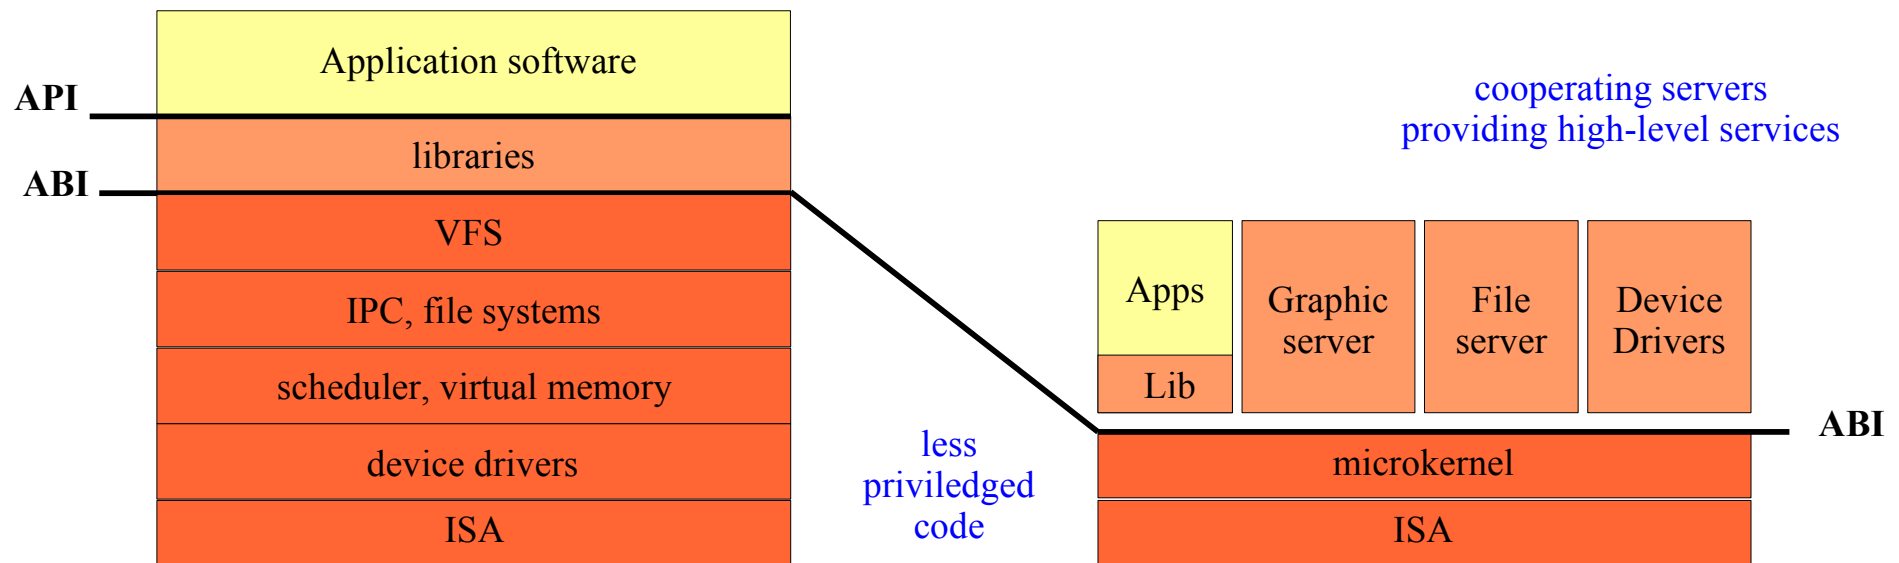
So now...
why do we have full network support
inside the kernel?

Operating System Architecture

19

- **Microkernels**

- Minimal kernel (TCB) with only few and low-level mechanisms
- Promote a modular design of cooperative user-level servers
- Cooperation relies on Inter-Process Communications (IPC)



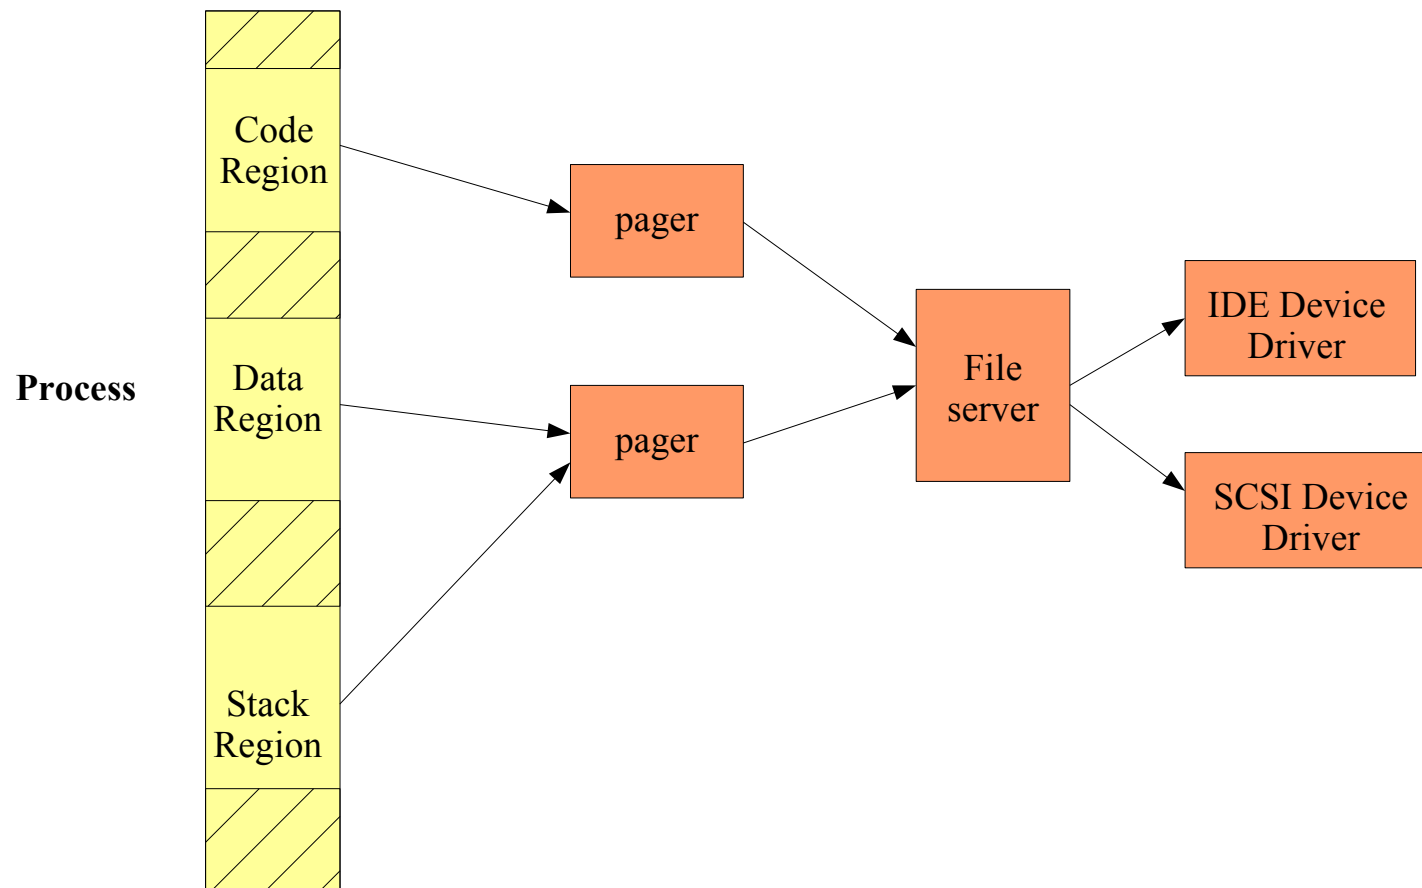
- **Processes as address spaces**
 - A virtualized memory, isolated by hardware
 - Mapping memory regions, managed by user-level pagers
- **Threads in address spaces**
 - Regular concept of threads (intra-process stacks of frames)
 - Microkernel is responsible for scheduling threads
 - Exceptions (traps) as IPCs
- **Inter-Process Communication (IPC)**
 - Message-based mechanism for inter-process communication
 - Sole medium of cooperation: send and receive messages

- Capabilities
 - Tokens representing privileges
 - The *right* to execute an action on an entity
 - Verified at runtime
- Different microkernels use capabilities
 - KeyKOS ('85), Mach ('87), EROS ('99)
 - OKL4 V2.1 ('08): first cap-based L4 kernel
 - SeL4 ('09): first proven microkernel

Microkernel – Processes

22

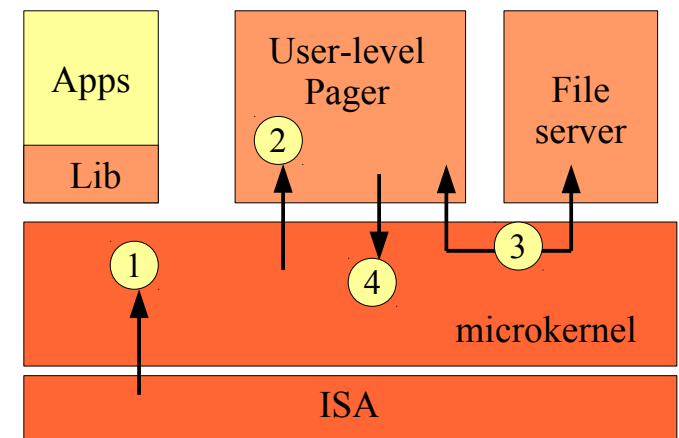
- Processes as address spaces
 - A virtualized memory, isolated by hardware
 - Mapping memory regions, managed by user-level pagers




User-Level Pagers

23

- Principle
 - The microkernel manages the page table and/or TLB
 - Relies on an external user-level server for paging
- Overview
 - (1) a page fault occurs
 - (2) Upcalls (IPC) the corresponding external pager
 - (3) External pager finds the page on disk
 - (4) Provide the page to the kernel



- Inter-Process Communications
 - Between end points (either ports or threads)
 - Payload as both in-line data and out-of-line data
 - One receiver, multiple senders
 - Blocking or non-blocking operations
 - Receiving on multiple ports with one thread is usually possible
- IPC Broker
 - How do you get a remote port?




- *Regular threads*
 - Execution flow, local to a process
 - Usually associated capabilities
 - Usually execute at a certain priority level
 - Uncaught exceptions redirected as messages to a handler end-point
 - Often has thread-local variables
- *Regular scheduling*
 - Microkernels embeds one or more scheduling policies
 - Often based on priorities and time slices
 - Executing thread can yield
 - Executing thread can send or receive messages
 - Synchronization mechanisms

Microkernels – Threads

26


- Discussing regular I/O
 - Application reads or writes a file or a socket
 - Translates to an IPC to the file server
- File Server
 - Manages the file cache (blocks are available or not)
 - IPC to the device driver to read missing blocks
- Device driver
 - Issues the I/O request to the device
 - Using memory-mapped I/O ports
 - Configured interrupts




External Device Drivers

27

- Interrupt handling
 - The microkernel captures the interrupts (1)
 - But it does not handle the interrupts
 - It forwards it to the concerned driver
 - Generating an IPC (2)
- The device driver handles the interrupt (3)
 - Potentially dialoging with the device through I/O ports
 - Interrupt may indicate the end of the DMA transfer



- Interrupt handling (cont...)
 - The device driver notifies the file server (4)
 - The file server caches the block and notifies the application (5)
 - The library unblocks the calling thread, the data being available



Warm-up discussion:

What are the differences between these three?

- a specification
- a design
- an implementation

Let's get to work:

- 1) What is the difference between a syscall and an IPC, if any?
- 2) Why do you need a service broker? Design one.
- 3) How do you create a process?
- 4) How do you map/unmap regions in a process?
- 5) How do you wrap the MMU into an IPC protocol with a pager?
- 6) How do you design malloc/free?
- 7) How do you create threads?
- 8) How do you manage thread stacks?
- 9) Specify send and receive syscalls
 - format of the messages, fixed or variable?
 - memory management, avoiding copies as much as possible
 - does the sender block if the receiver isn't ready?
 - does the receiver block if there is no message
 - IPC end points? threads? Mailboxes or ports ?
- 10) Design send and receive syscalls

Bonus: How do you bootstrap?

Discussing Microkernels

30

- Some spectacular “failures”
 - Most notorious: IBM Workspace OS
 - Also the GNU Hurd
- Recently spectacular “successes”
 - Apple Mac-OS
 - GreenHills (world leader in RTOS)
 - Nicta (1.5 billions smart phones use L4)

- Unify IBM's operating systems (DOS, OS/2, AIX, OS/400, 390, ...)
 - Cost saving but also multiple personalities on the same hardware
 - Based on Mach 3.0 from Carnegie Mellon University (CMU)
 - Across a wide range of hardware environments
 - PDAs, desktops, massively parallel machines
 - With concurrent OS personalities
 - Sharing personality neutral services (PSNs)
 - Applications could use services from multiple OSes
- Very ambitious project (budget and manpower)
 - Plagued by internal problems
 - Too much focus on portability, not enough on personalities
 - Microkernel performance not mentioned as one of the main issues

- Nevertheless, disappointing first generation
 - By 1991, an IPC overhead bottomed at 115 μ s on a 486-DX50
 - Conventional Unix system call was at roughly 20 μ s
 - About 10 times faster (syscall = 2 IPCs)
 - But it is hard to discuss and appreciate micro-benchmark numbers
- Overall results
 - Some applications performed as well or slightly better on microkernels
 - Some showed peak degradation of 50% to 66%
 - Clearly depends on how much applications use system calls

- Analysis
 - At least 73% of the measured penalty was attributed to IPC overhead
 - 10% was attributed to multi-processor provision in microkernel that was not present in the uniprocessor Unix measured
 - 17% were unattributed
- Conclusions
 - IPC-based architecture was flawed from a performance perspective
 - Mach and Chorus microkernels reintegrated performance-crucial servers
 - Both had limited commercial successes
 - IPC performance is believed to be the main rationale for the commercial failure of microkernels

- Which applications were used?
 - Regular applications as sed, egrep, yacc, gcc or compress
 - Are they really representative?
- What would be the case for Open Office?
 - More CPU intensive and less I/O bound
 - For instance, graphics are already IPC-based (with X11)
- Was 50-60% degradation a real showstopper?
 - It was considered as one, by researchers themselves
 - Industry was not to accept something researchers found a failure
 - What about Moore's law
 - Hardware was twice faster every year

Discussing Commercial Failure

35

- The overhead can't justify commercial failure
 - Did you notice that Java failed in its early years because it was slow?
 - What about newer versions of Microsoft Windows?
- Really missing a killer reason for adoption
 - Flexibility, modularity, and freedom of choice are not obvious market growth opportunity
 - Forgot the natural resistance to change bottom layers of complex critical software stacks
- Who was to pay for porting existing OSes on microkernels?
- **Who was to rewrite all device drivers?**

Discussing Commercial Failure

36

- We have a chicken and egg problem
 - New applications requiring a specific OS platform is a commercial No-No
 - Without new applications, microkernel were not seen as an advantage
- Traditional Unix systems evolved proposing new system calls
 - Memory mapped files
 - Pinning pages in memory
 - Efficient signals for virtual memory protection violations
- Middleware phenomenon versus multiple personalities
 - We saw cross-porting of much libraries and file-systems
 - We saw cross-porting of applications
 - We also saw distributed integration
- Hypervisor phenomenon
 - Multiplexing operating systems

Stepping Back...

37

- What about our numbers and our understanding of them?
 - Numbers clearly showed that IPC was the crux of the problem
 - At about 100 μ s on a 486-DX50, it seemed an unacceptable overhead
 - The conclusion was that approach was flawed and doomed
- Was that correct? Well... not really
 - Naive and uninterpreted measurements can be misleading
 - Performance measurements on complex systems are not trivial
- Reality
 - It was the implementation of the kernels that was at fault
 - The code was borrowed from monolithic kernel
 - It was improved incrementally (usually considered a good idea)
 - But the approach proved to be at the heart of the problem
 - Reducing a large kernel does not produce an efficient micro-kernel

- Case study: L4 microkernel
 - Developed at GMD in 1995
 - Provides a minimal set of abstractions
 - Like the first generation
 - But argues non-portable implementations
 - For real performance, implementations must be processor dependent
 - So writing microkernel code is not about writing portable code
 - Even between 486 and Pentium!

Jochen Liedtke, *Towards Real Microkernels*
Communications of the ACM, Vol 39, No 9, 1996

Jochen Liedtke, *Improving IPC by Kernel Design*,
SIGOPS'93

Introducing L4 Microkernel

40

- L4 Abstractions
 - Task: an address space with threads
 - Inter-Process Communications (IPC)
- L4 Design
 - Only 7 system calls while Mach has 140 system calls
 - Footprint of 12KB while Mach footprint is 300KB
- L4 Performance (on 486-DX50)
 - 8-byte IPC
 - **5 μ s** for L4 against 100 μ s for Mach
 - L4 IPC (2*5 μ s) is twice as fast as Unix system call
 - 512-byte IPC
 - **18 μ s** for L4 against 172 μ s for Mach
 - Only twice slower than a Unix system call (2*18 μ s versus 20 μ s)

From L3 to L4 – Starting Fresh

41

- Principles and methods
 - IPC performance is the master design issue
 - All design decisions require a performance decision
 - If something performs poorly, look for new techniques
 - Synergetic effects have to be taken into considerations
 - The design has to cover all levels from architecture down to coding
 - The design has to be made on a concrete basis (hardware capabilities)
 - The design has to aim at a concrete performance goal
 - Only accept a concept in the kernel if it cannot be implemented above


- Concrete design
 - 486-based hardware family
- Intel 486-DX50
 - Clock rate 50MHz
 - On-chip cache of 8KB
 - Typical instruction takes 1 or 2 cycles, assuming cache hits
 - The MMU translates 32-bit virtual addresses
 - Using 4KB pages
 - With read-write protections
 - TLB has 32 entries
 - TLB is flushed on switching address space

- Simple scenario – Zero-byte message
 - Two threads A and B in two address spaces
 - Sending a zero-byte message, from thread A to thread B already waiting
- Rough estimate
 - Down to 20 instructions
 - Ignoring parameter passing and testing
 - These 20 instructions are about 127 cycles
 - Out of which 107 cycles are consumed to enter/leave kernel mode
 - 486 MMU is flushed on address space switch
 - So we need to add at least 5 TLB misses
 - One TLB miss is about 9 cycles
 - So this means a lower bound of 172 cycles (about 3.5 μ s)
 - Not realistic, but a lower bound
 - Since we assumed no write or code prefetch delays or cache misses


- A realistic target was 350 cycles, about $7\mu\text{s}$
 - Reach **250 cycles, $5\mu\text{s}$**
 - **$T=5\mu\text{s}$** will be used as a basic unit for discussing performance
- How was it done?

- Reducing system calls
 - System calls are expensive
 - 107 cycles for entering and leaving kernel, about 2 μ s
 - So about 40% T
 - Traditional RPC
 - Two send and receive IPC calls on both client and server side
 - Results in 4 enter-leave kernel-mode instructions
 - Combining IPCs
 - IPC call combines sending a message and receiving the reply
 - IPC call combines replying and receiving next message
 - We are down to only 2 enter-leave instructions

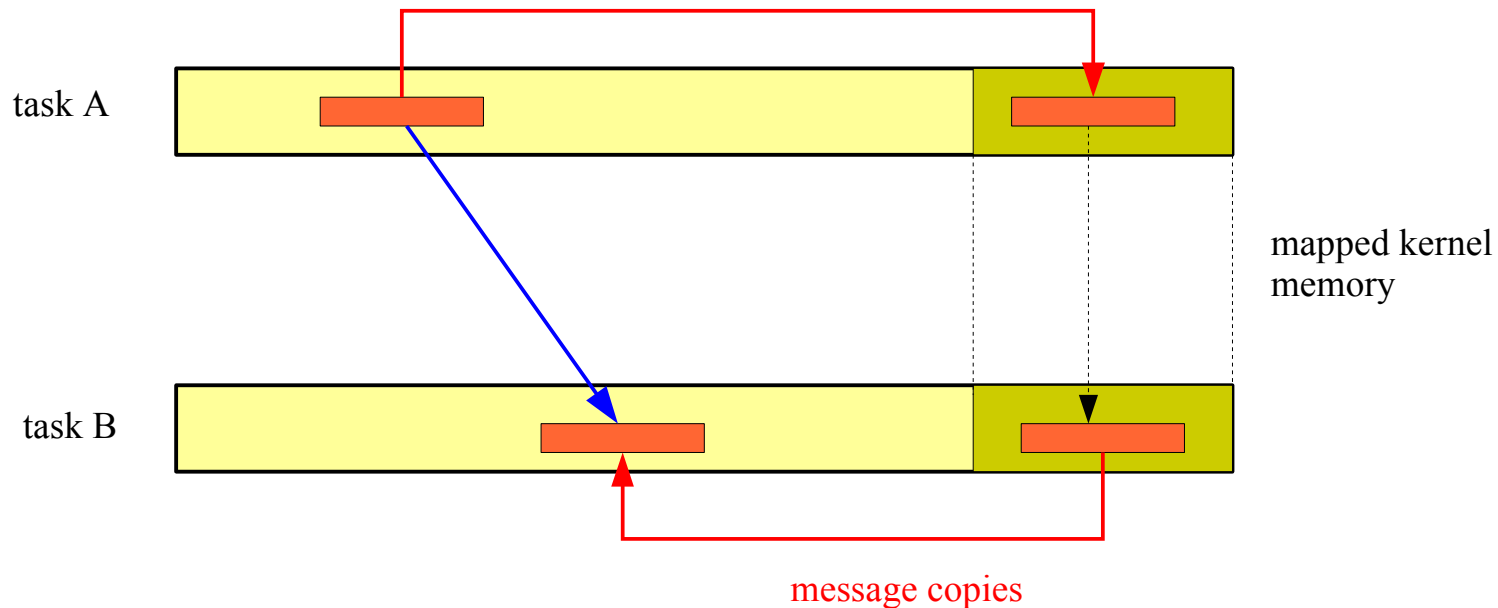
- Overhead of system calls and address space switch
 - About 60% T, which suggests to have complex message à la Mach
 - Sending larger messages reduce the relative overhead
- Message format
 - With mandatory and optional strings
 - And optional memory objects
 - Provides a gather-scatter approach



- Optional memory objects
 - Passed by playing with address space mappings
- Optional strings
 - Passed through copies between address spaces
 - But traditional application-level copies are avoided
 - Message is constructed using the address of language variables



- Traditional design
 - Costly copies, from sender to kernel and from kernel to receiver
- Overhead
 - For n-byte message, about $20 + 0.75n$ cycles
 - 8-byte message is about 26 cycles, about $0.5\mu\text{s}$, about 10% T
 - 512-byte message is about 404 cycles, about $8\mu\text{s}$, about 160% T



IPC Performance – Lightweight RPC

49

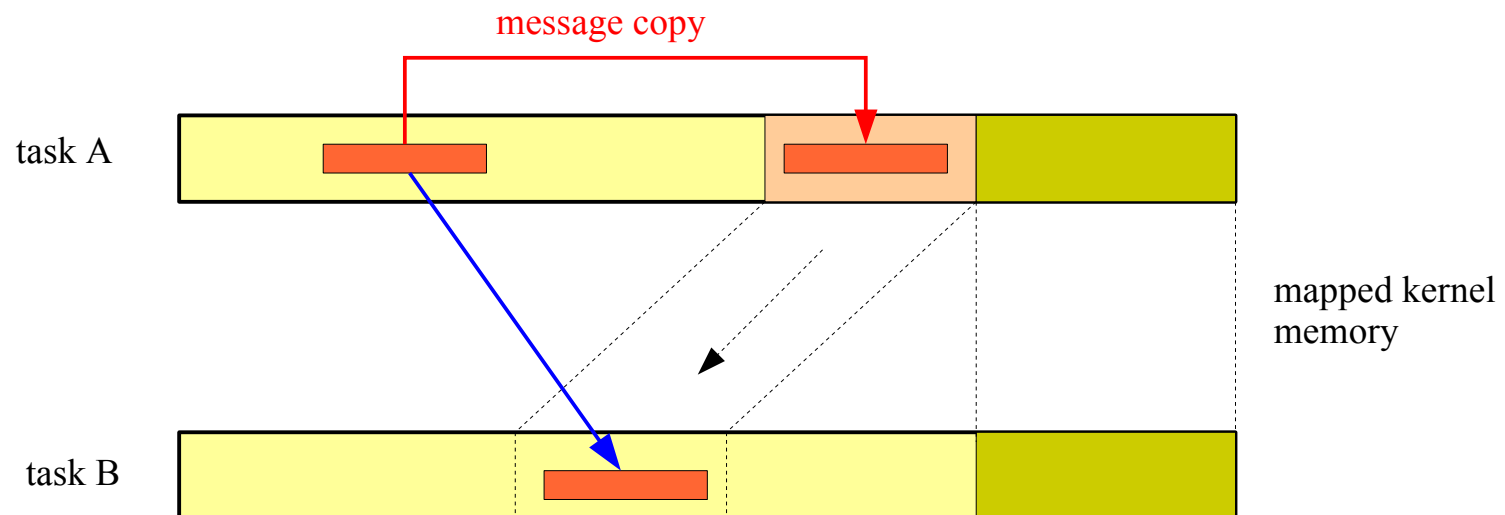
- Lightweight RPC
 - Based on user-level shared buffers
 - Requires an explicit open of connections to allocate buffers
 - One copy in sender address space from message to the shared buffer
- Does not scale well
 - For server tasks with many clients
 - Shared buffers with large messages may become a scarce resource
- Security breach
 - Receiver can't ensure the legality of the message
 - Since the sender may change it after the receive
 - Constitute hidden channels without proper IPC controls

Bershad et al, Lightweight Remote Procedure Call,
ACM Symposium on Operating System Principles, Dec 1989, pp 102-113

IPC Performance – Single-Copy Design

50

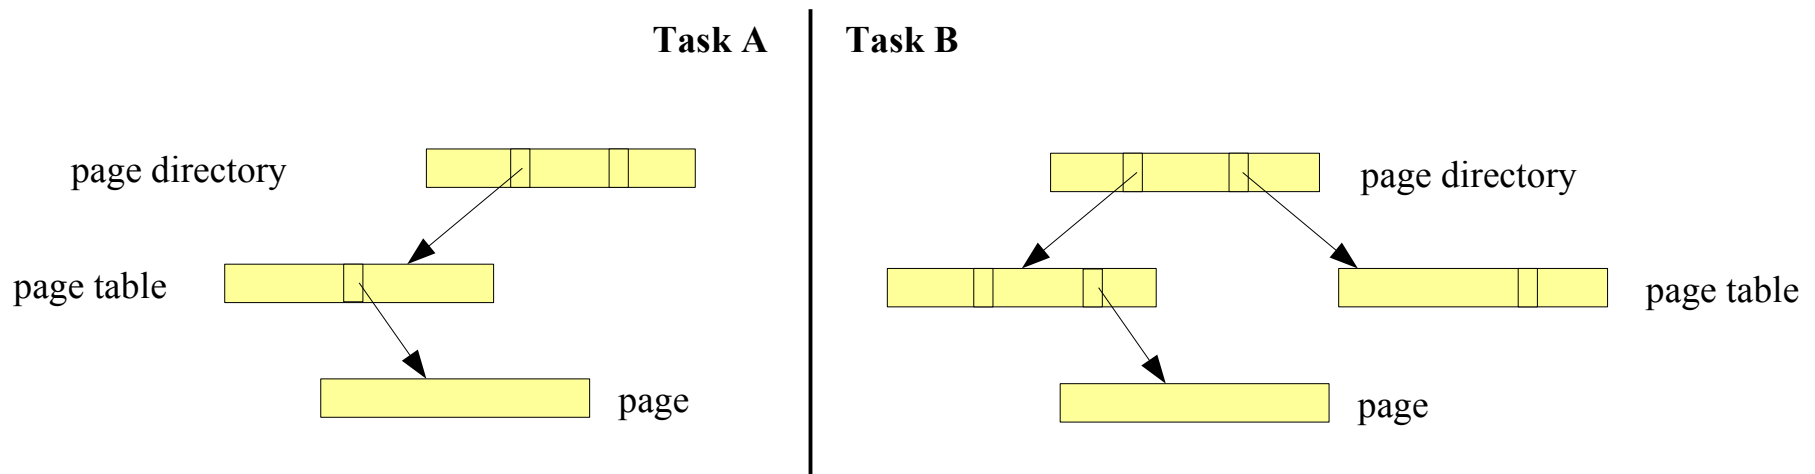
- Direct transfers by temporary mappings
 - Map the necessary virtual pages
 - Temporary mapping in task A, only accessible to the kernel
 - Copy only once the message
 - Unmap the temporary mapping



IPC Performance – Single-Copy Design

51

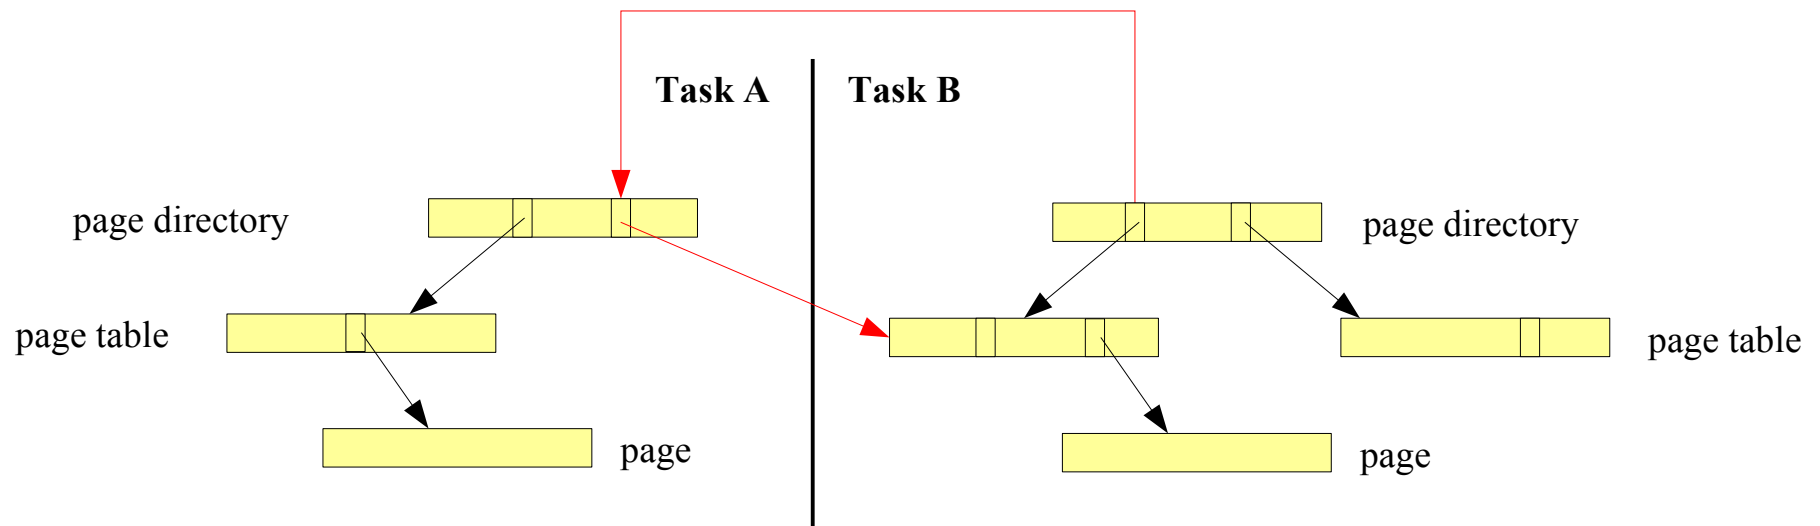
- Temporary mapping design
 - Two challenges
 - Temporary mapping must be fast
 - Design must work with multiple threads in the same address space
 - Performance on 486 hardware
 - 486 MMU uses an architected two-level page table
 - Page directory is the first level, hold 1024 entries
 - Page tables are the second level, each also holds 1024 entries to pages



IPC Performance – Single-Copy Design

52

- Temporary mapping design
 - Two-word-copy mapping on 486 hardware
 - Each page table corresponds to 4MB
 - To map aligned 4MB, we need one word copy between page directories
 - Messages are restricted to 255 strings and always map to 8MB region

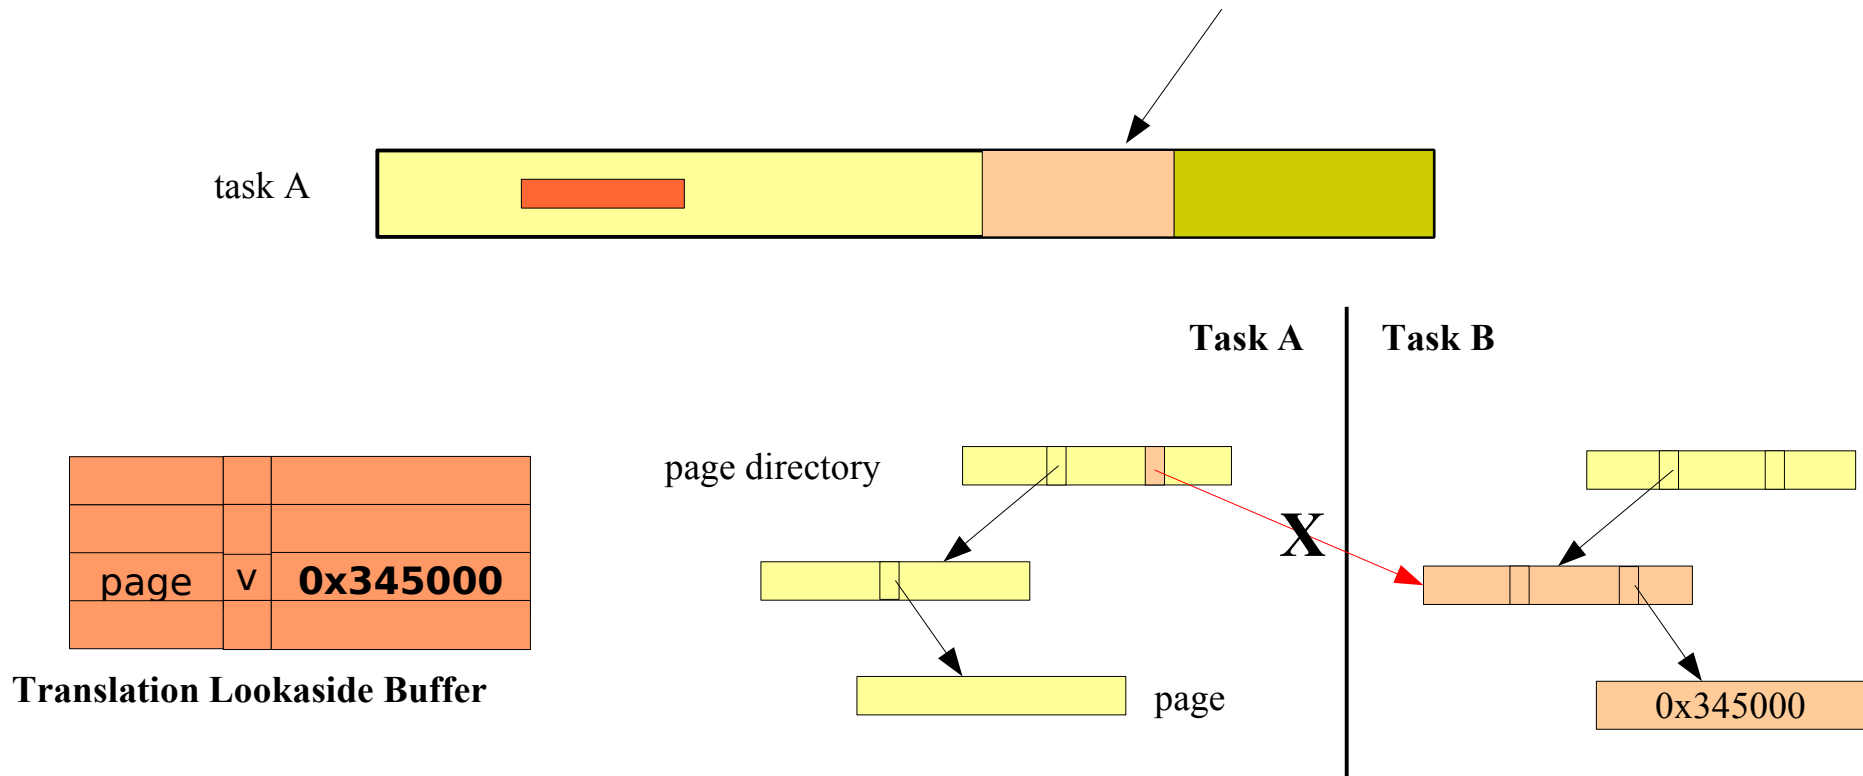


IPC Performance – Single-Copy Design

53

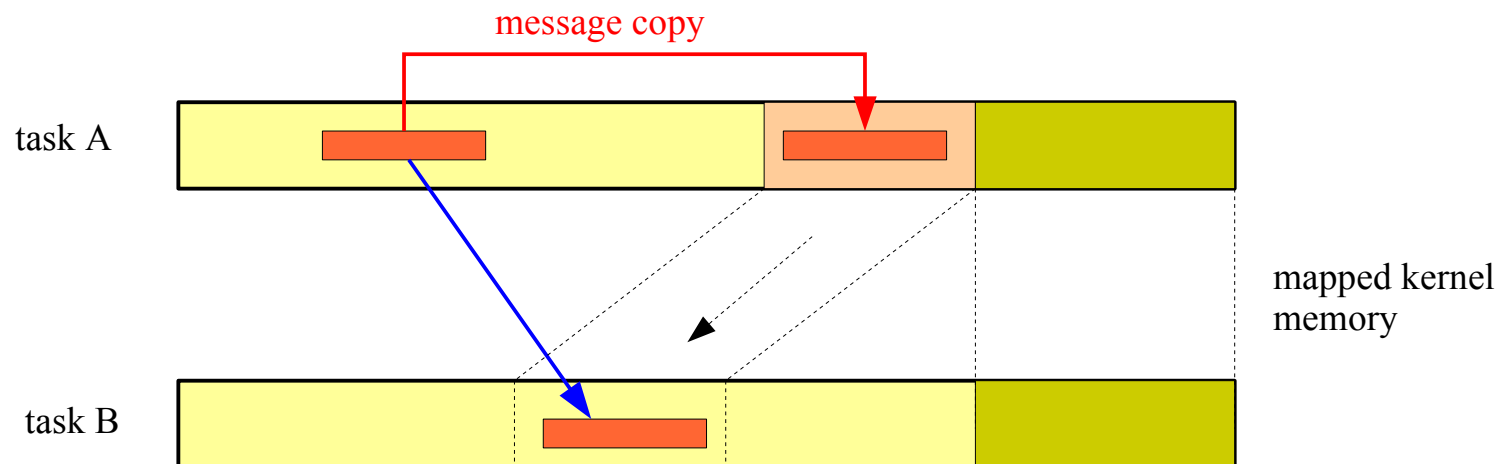
- Removing the temporary mapping
 - Cutting the entry in the page directory is not enough
 - We need to ensure that we have a *IPC-clean* TLB:

An **IPC-clean TLB** has no entries related to any *IPC-mapped region*



- Ensuring an IPC-clean TLB
 - Flushing is too expensive
 - Flushing an entire TLB induces a too great overhead
 - Flushing individual pages for 4MB regions is too slow
- A solution with one thread per address space
 - When starting a task
 - The TLB is *IPC-clean* since it is empty
 - When switching to a thread in a task
 - The TLB is flushed on address space switch
 - Therefore the TLB *IPC-clean*
 - What about during the IPC?

- During the IPC:
 - (1) map the IPC region
 - (2) do the message copy
 - (3) unmap IPC region
 - (4) switch to receiver thread
- In step (4), the TLB is IPC-clean
 - We switch address space (since we have one thread per task)
 - So we flushed the TLB, so it is IPC-clean



- Discussing multiple threads
 - One 4MB region per thread may be difficult to get (only 1024 entries)
 - Use only one IPC-mapped region, shared by all threads
- First design
 - Synchronization
 - Only one thread can use the IPC-mapped region
 - Mutual exclusion only during the send part of the IPC
 - Mapping management
 - Restore the correct mapping upon thread switching
 - Since each thread has potentially its own mapping for the IPC-mapped region

- Scheduler modifications
 - When scheduling between threads
 - Always invalidate the IPC region in the page table (no access allowed)
 - Between threads in different address spaces
 - Flush the TLB (as before)
- Pager modification
 - Upon any page fault in the IPC region
 - Restores the correct mapping for the current thread
 - At page fault, we know the address and faulting thread
 - Since a given thread may be in only one IPC at a time
 - We know the IPC region to restore

- Key design points
 - Used Rendez-Vous
 - Combined send and receive (half the number of system calls)
 - Only one copy through temporary virtual memory mapping
 - Required to change the scheduler to ensure IPC-clean TLB
 - Switch to receiver upon rendez-vous
 - Thread-aware temporary mapping, with pager support
- Now what?
 - Managing timeouts and wakeups...

IPC Performance – Timeouts & Wakeups

59

- Timeouts may be specified on each IPC
 - A timeout of τ means that the thread is awakened if message transfer has not started τ ms after invoking the operation
- Timeouts require a wakeup queue
 - For each IPC, insert the thread in the wakeup queue
 - Upon IPC completion, remove the thread from the wakeup queue
- Timer interrupt is expensive
 - Must scan the wakeup queue
 - Test for the expiration of each timeout value
 - For expired timeouts, awake the corresponding thread
 - Remove it from the wakeup queue
 - Insert it in the ready queue

IPC Performance – Timeouts & Wakeups

60

- Fast wakeup queue operations
 - Simple and fast design is an array
 - Indexed by the thread number
 - Holding the wakeup time, computed from timeout values
 - Insert/remove from the wakeup queue is then extremely fast
 - Just an array store (with zero or a wakeup time)
- But the timeout management is costly
 - Timer interrupt handler must scan the array
 - For 16K entries
 - This would require 2ms for each timer interrupt
 - On-chip cache would be flooded
 - Not acceptable
 - Why so many threads?
 - Remember, blocking IPC tend to require many threads

IPC Performance – Timeouts & Wakeups

61

- A better-balanced design
 - Use N lists for the wakeup queue
 - A thread with a timeout τ is inserted in the $(\tau \bmod N)$ list
- Timer interrupt handler
 - For k threads, on average, scans k/n entries
- Optimization to reduce scanning time further
 - Divide each of the N wakeup queues
 - Create a short-timeout and a long-timeout wakeup queue
 - Insert in the short-timeout queue threads whose timeout approaches

IPC Performance – Timeouts & Wakeups

62

- For $n=8$ lists
 - A wakeup granularity of 4 ms (check timeout list 250 times per second)
 - With 400 threads
 - We inspect at most $250 \times 400 / 8 = 12500$ inspected entries per second
 - We spend less than **1%** of processor cycles in scheduling
 - At this rate of 12500 entries inspected per second
 - This would require at least 12500 ipc per second
 - Would mean at least 6% of the CPU time used in pure IPC
 - In practice
 - Since 75% of the IPC have concrete timeout values others than 0 and ∞
 - This scenario would require 50,000 IPC/sec
 - At 50000 IPC/sec, IPCs would take **25%** of the processor time
 - Scheduling overhead is therefore about 4% of that total IPC time
 - $4\% = 1/25$ (the 1% and 25% processor cycles above)

- Long Timeouts and wakeups
 - Even though timeouts may be limited to 2^{31} ms, about 24 days
 - Corresponding wakeup times will be longer
 - 64bit quantities are too expensive
 - Using too many registers on a 32bit processor
- Use base+offset to represent a wakeup time
 - Manage time offsets always smaller than 2^{24} ms
 - Every 4.5 hours (2^{24} ms), the kernel updates the base
- Wakeup lists managed with a base and an offset
 - Split wakeups between those in the current window and later wakeups
 - On the time interrupt, only scan wakeups in the current window (2^{24} ms)


- Key design points
 - Used Rendez-Vous
 - Combined send and receive (half the number of system calls)
 - Only one copy through temporary virtual memory mapping
 - Required to change the scheduler to ensure IPC-clean TLB
 - Switch to receiver upon rendez-vous
 - Thread-aware temporary mapping, with pager support
 - Efficient management of wakeup queues
- Now what?
 - Address spaces...

- History
 - First generation microkernels had limited flexibility
 - Provided only the concept of external pagers, with mapped regions
 - Replacement policy were not accessible
 - L4 approach
 - Only mechanisms in the kernel
 - Powerful enough to build policies outside the kernel
- Core ideas
 - Recursive construction of address spaces through mappings
 - An *initial space* represents the physical memory
 - Address space owners
 - Threads executing in that address space
 - Only three operations by owners
 - **Grant** or **Map** a page to another address space
 - **Demap** a page from other address spaces

L4 Address Spaces

66


- **Map** IPC between owner threads
 - `send('map',virtual @)`
 - `receive('map',virtual @)`
- **Semantics**
 - Applies to a page
 - Sender retains the page in its own address space
 - Receiver accepts the page mapping, chooses the address



L4 Address Spaces

67

- **Grant** IPC between owner threads
 - `send('grant',virtual @)`
 - `receive('grant',virtual @)`
- **Semantics**
 - Applies to a page
 - Sender **does not retain** the page in its own address space
 - Receiver accepts the page mapping, chooses the address



- Memory mappings

$$\sigma : V \rightarrow R \cup \{\emptyset\}$$

where V is the set of virtual pages

R is the set of available real pages (physical)

\emptyset is the nil page that cannot be accessed

Initial mapping representing the physical memory

$$\sigma_0(v) = v$$


Map operation: $\sigma_j(v') = \sigma_i(v)$

Grant operation: $\sigma_j(v') = \sigma_i(v)$
 $\sigma_i(v) = \emptyset$

L4 Address Spaces

69


- A hierarchy of address spaces
 - Each owner may map, to other address spaces, pages that it owns



L4 Address Spaces

70

- Demap operation
 - An owner may demap one of its page
 - It retains the mapping of that page
 - For all *higher* address spaces that obtained that page
 - Directly or not from that owner
 - They loose that page



- L4 address space model

$$\sigma : V \rightarrow R \cup \{\emptyset\}$$

where V is the set of virtual pages

R is the set of available real pages (physical)

\emptyset is the nil page that cannot be accessed

Initial mapping representing the physical memory

$$\sigma_0(v) = v$$


Map operation: $\sigma_j(v') = \sigma_i(v)$

Grant operation: $\sigma_j(v') = \sigma_i(v)$
 $\sigma_i(v) = \emptyset$


Demap operation: $\forall j, v'$ such that $\sigma_j(v') = \sigma_i(v)$
 $\sigma_j(v') = \emptyset$

- L4 tasks
 - Single address space
 - A flat virtual address space (no segments)
 - One or more threads
 - Threads are called owners of the address space
 - Notation: $\theta_i(\sigma_j)$ for a thread owning an address space
 - An address space is created with one owner
 - A pager
 - A thread to receive page fault traps
 - Traps are translated to IPC in L4

task

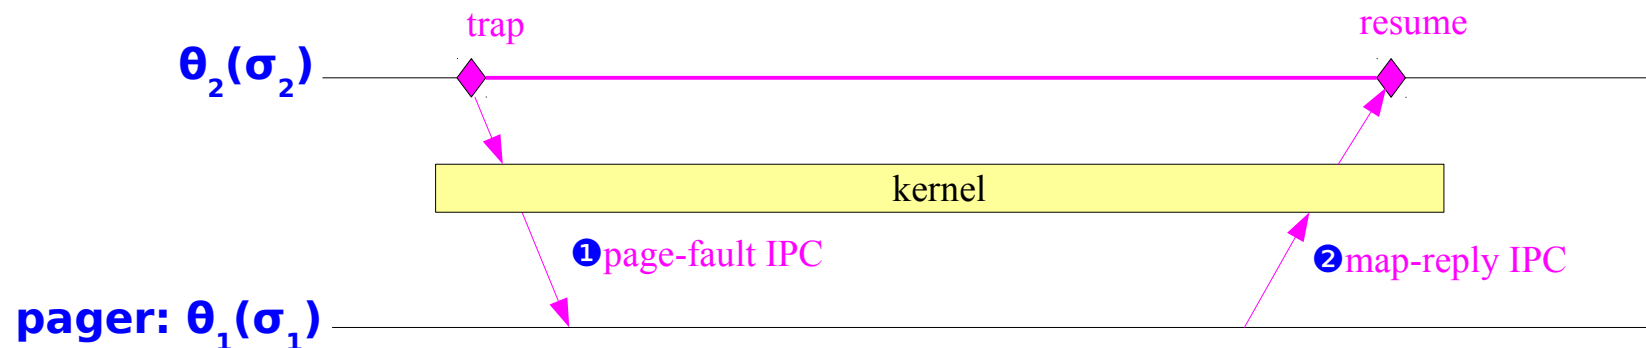



- Illustrating address spaces
 - Initial pager task
 - Simple mapping scheme
 - Maps each page only once, never demaps pages
 - Virtual memory pager
 - Requests a number of physical pages
 - That it manages as a cache

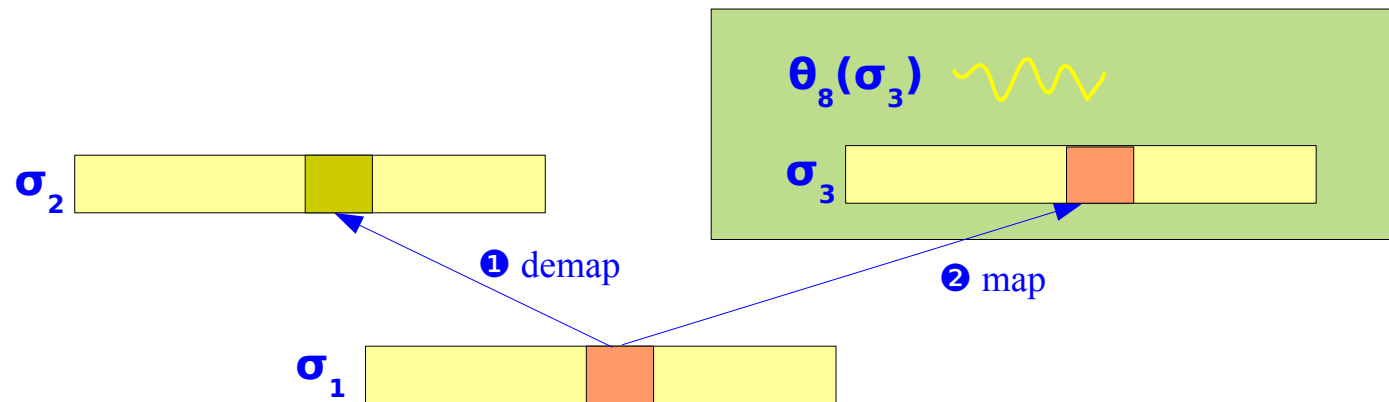


- Illustrating page faults

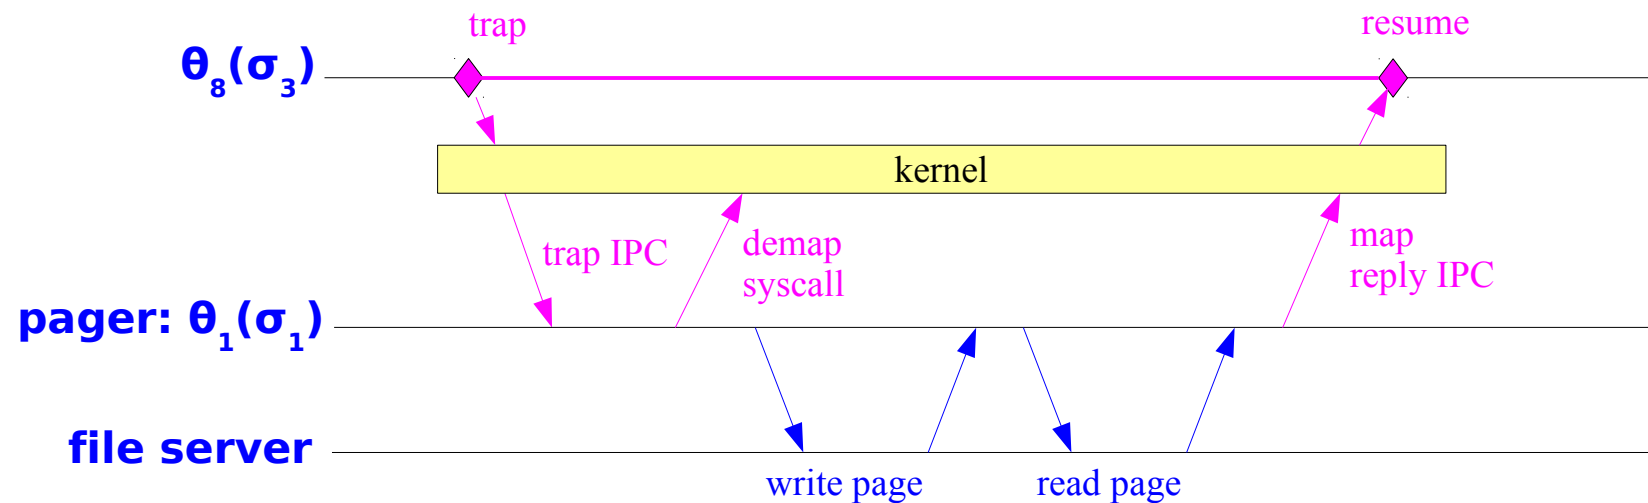
- $\theta_2(\sigma_2)$ traps
 - Trap is translated into an IPC to the pager
 - Pager on thread $\theta_1(\sigma_1)$
- Pager handles the page faults
 - Selects a free page in the cache
 - Loads the content of that page
 - Maps that page



- Illustrating page replacement
 - $\theta_8(\sigma_3)$ traps
 - Trap is translated into an IPC to the pager
 - No more physical pages are available
 - Page replacement
 - Demap one of its managed page
 - Reuse that page to map it



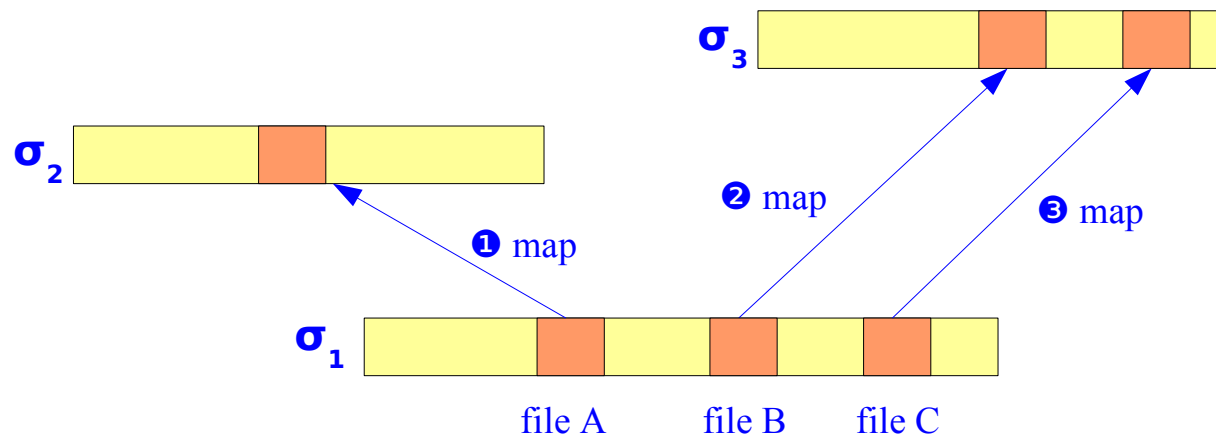
- Illustrating swapping
 - $\theta_8(\sigma_3)$ traps
 - Trap is translated into an IPC to the pager
 - Pager handles the page faults
 - Write out the cached page and then demap it
 - Read in the new content of the page
 - Maps the page



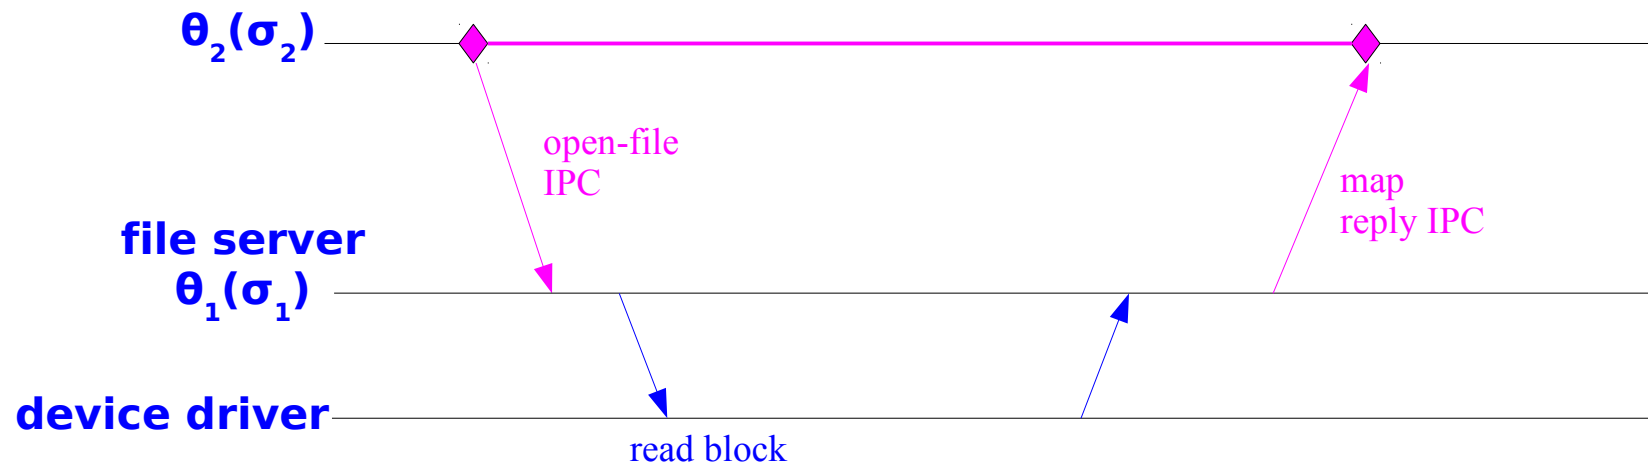
Microkernel Architecture

77

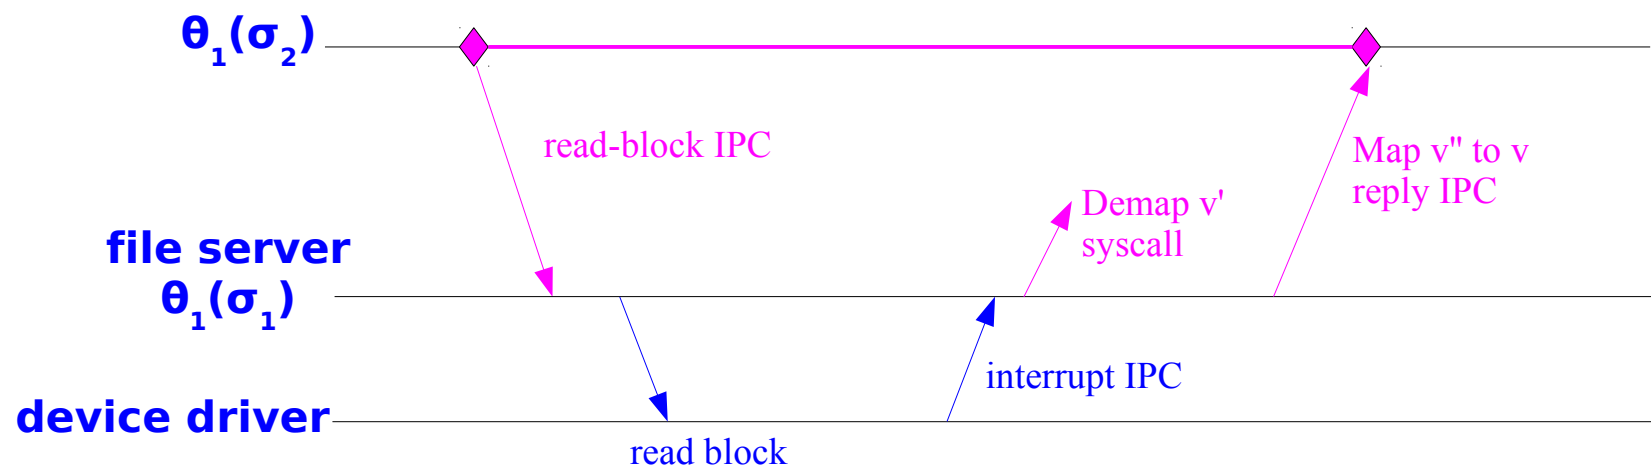
- Illustrating a simple file server
 - Manages a cache of physical pages
 - Use them as read-write buffers
 - For each client
 - Opened files and current position (I/O block)
 - Current block is mapped in client address space



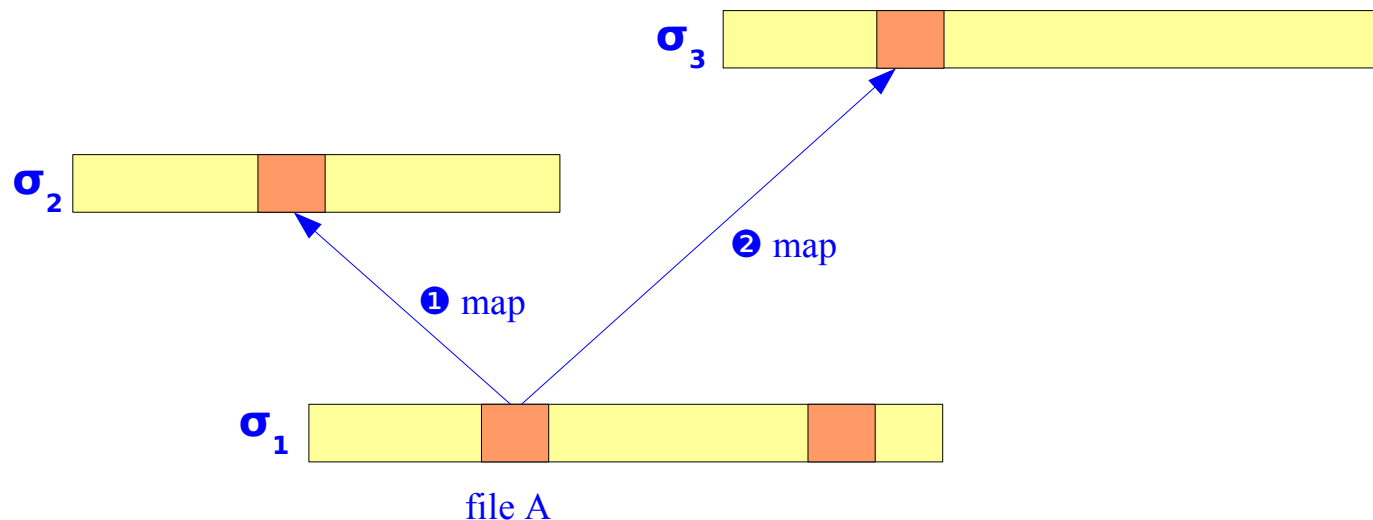
- Illustrating a simple file server
 - Opening a file (client)
 - Client allocates a buffer (a page v in virtual memory)
 - Calls the file server with a open-request IPC
 - Waits for a reply IPC, a map request: $\sigma_2(v) = \sigma_1(v')$
 - File server
 - Does an internal open of the file
 - Allocates a page v' in its cache and reads in the first block
 - Replies a map IPC, $\sigma_2(v) = \sigma_1(v')$



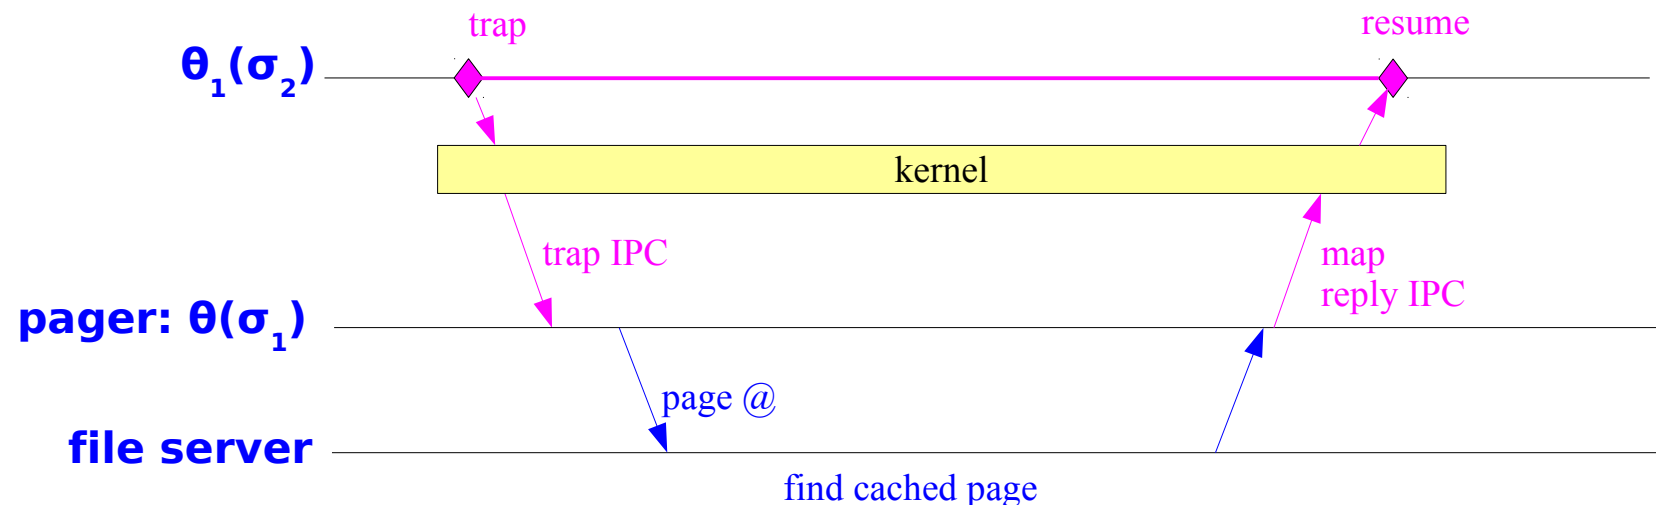
- Illustrating a file server
 - Read a block, given a block offset
 - Calls the file server with a read-block-at-offset IPC
 - Waits for a reply IPC, a map request **at the same client buffer v**
 - From $\sigma_2(v) = \sigma_1(v')$ to $\sigma_2(v) = \sigma_1(v'')$
 - File server
 - Allocates a new page v'' in its cache and initiate I/O
 - Demaps the previous page v' that held the previous current block
 - Replies a map IPC for the new page holding the new current block



- Illustrating a file server
 - Concurrently opened files
 - Current blocks may be mapped in different clients
 - Always true for concurrently-opened files smaller than 4KB on most systems
 - Demap challenge
 - Demap system call unmaps $\sigma_i(v)$ from all clients
 - $\forall j, v'$ such that $\sigma_j(v') = \sigma_i(v) : \sigma_j(v') = \emptyset$
 - Page faults will occur, how do we re-establish the mapping?



- Illustrating a file server
 - A possible solution (others probably exist)
 - Specific address space structure
 - Reserve a file-system buffer region
 - File-system-aware pager
 - Handles the page fault if not in the file-system buffer region
 - Otherwise forwards the page trap to the file server
 - File server finds its corresponding page buffer and makes the map reply



- Architectural debates
 - Tanenbaum/Torvalds debate part 1 (1992)
 - <http://www.oreilly.com/catalog/opensources/book/appa.html>
 - Tanenbaum/Torvalds debate part 2 (2006)
 - <http://www.cs.vu.nl/~ast/reliable-os/>
 - <http://www.coyotos.org/docs/misc/linus-rebuttal.html>
- On micro-kernels versus virtual machine monitors
 - S. Hand et al., Are virtual machine monitors microkernels done right?, Proceedings of the 10th Workshop on Hot Topics in Operating Systems (HotOS-X), 2005.
 - G. Heiser et al., Are virtual machine monitors microkernels done right? ACM Operating Systems Review, Volume 40, Issue 1, 2006.
 - T. Roscoe et al., Hype and virtue, Proceedings of the 11th Workshop on Hot Topics in Operating Systems (HotOS-XI), 2007.

- Linux kernel software engineering study
 - Stephen R. Schach, Bo Jin, David R. Wright, Gillian Z. Heller, and A. Jefferson Offutt. Maintainability of the Linux kernel. IEE Proceedings: Software, 149:18–23, 2002.
- Origins of microkernels
 - Per Brinch Hansen. The nucleus of a multiprogramming operating system. Communications of the ACM, 13:238–250, 1970.
 - W. Wulf, E. Cohen, W. Corwin, A. Jones, R. Levin, C. Pierson, and F. Pollack. HYDRA: The kernel of a multiprocessor operating system. Communications of the ACM, 17:337–345, 1974.

- Mach

- M. Acetta et al., Mach: A new kernel foundation for UNIX development. Proceedings of the 1986 Summer USENIX Technical Conference, 1986.
- David B. Golub, Randall W. Dean, Alessandro Forin, and Richard F. Rashid. UNIX as an application program. In Proc. 1990 Summer USENIX Annual Technical Conference, pages 87–95, June 1990.

- L4

- J. Liedtke, On μ -Kernel Construction. Proceedings of the 15th ACM Symposium on Operating System Principles (SOSP), 1995.
- J. Liedtke, Toward Real μ -kernels. Communications of the ACM, 39(9), pp. 70-77, 1996.
- Hermann Härtig, Michael Hohmuth, Jochen Liedtke, Sebastian Schönberg, and Jean Wolter. The performance of μ -kernel-based systems. In Proc. Symp. on Operating Systems Principles, pages 66–77, St. Malo, France, October 1997.
- Gerwin Klein et al. seL4: Formal verification of an OS kernel. Proceedings of the 22nd ACM Symposium on Operating Systems Principles, Big Sky, MT, USA, October, 2009